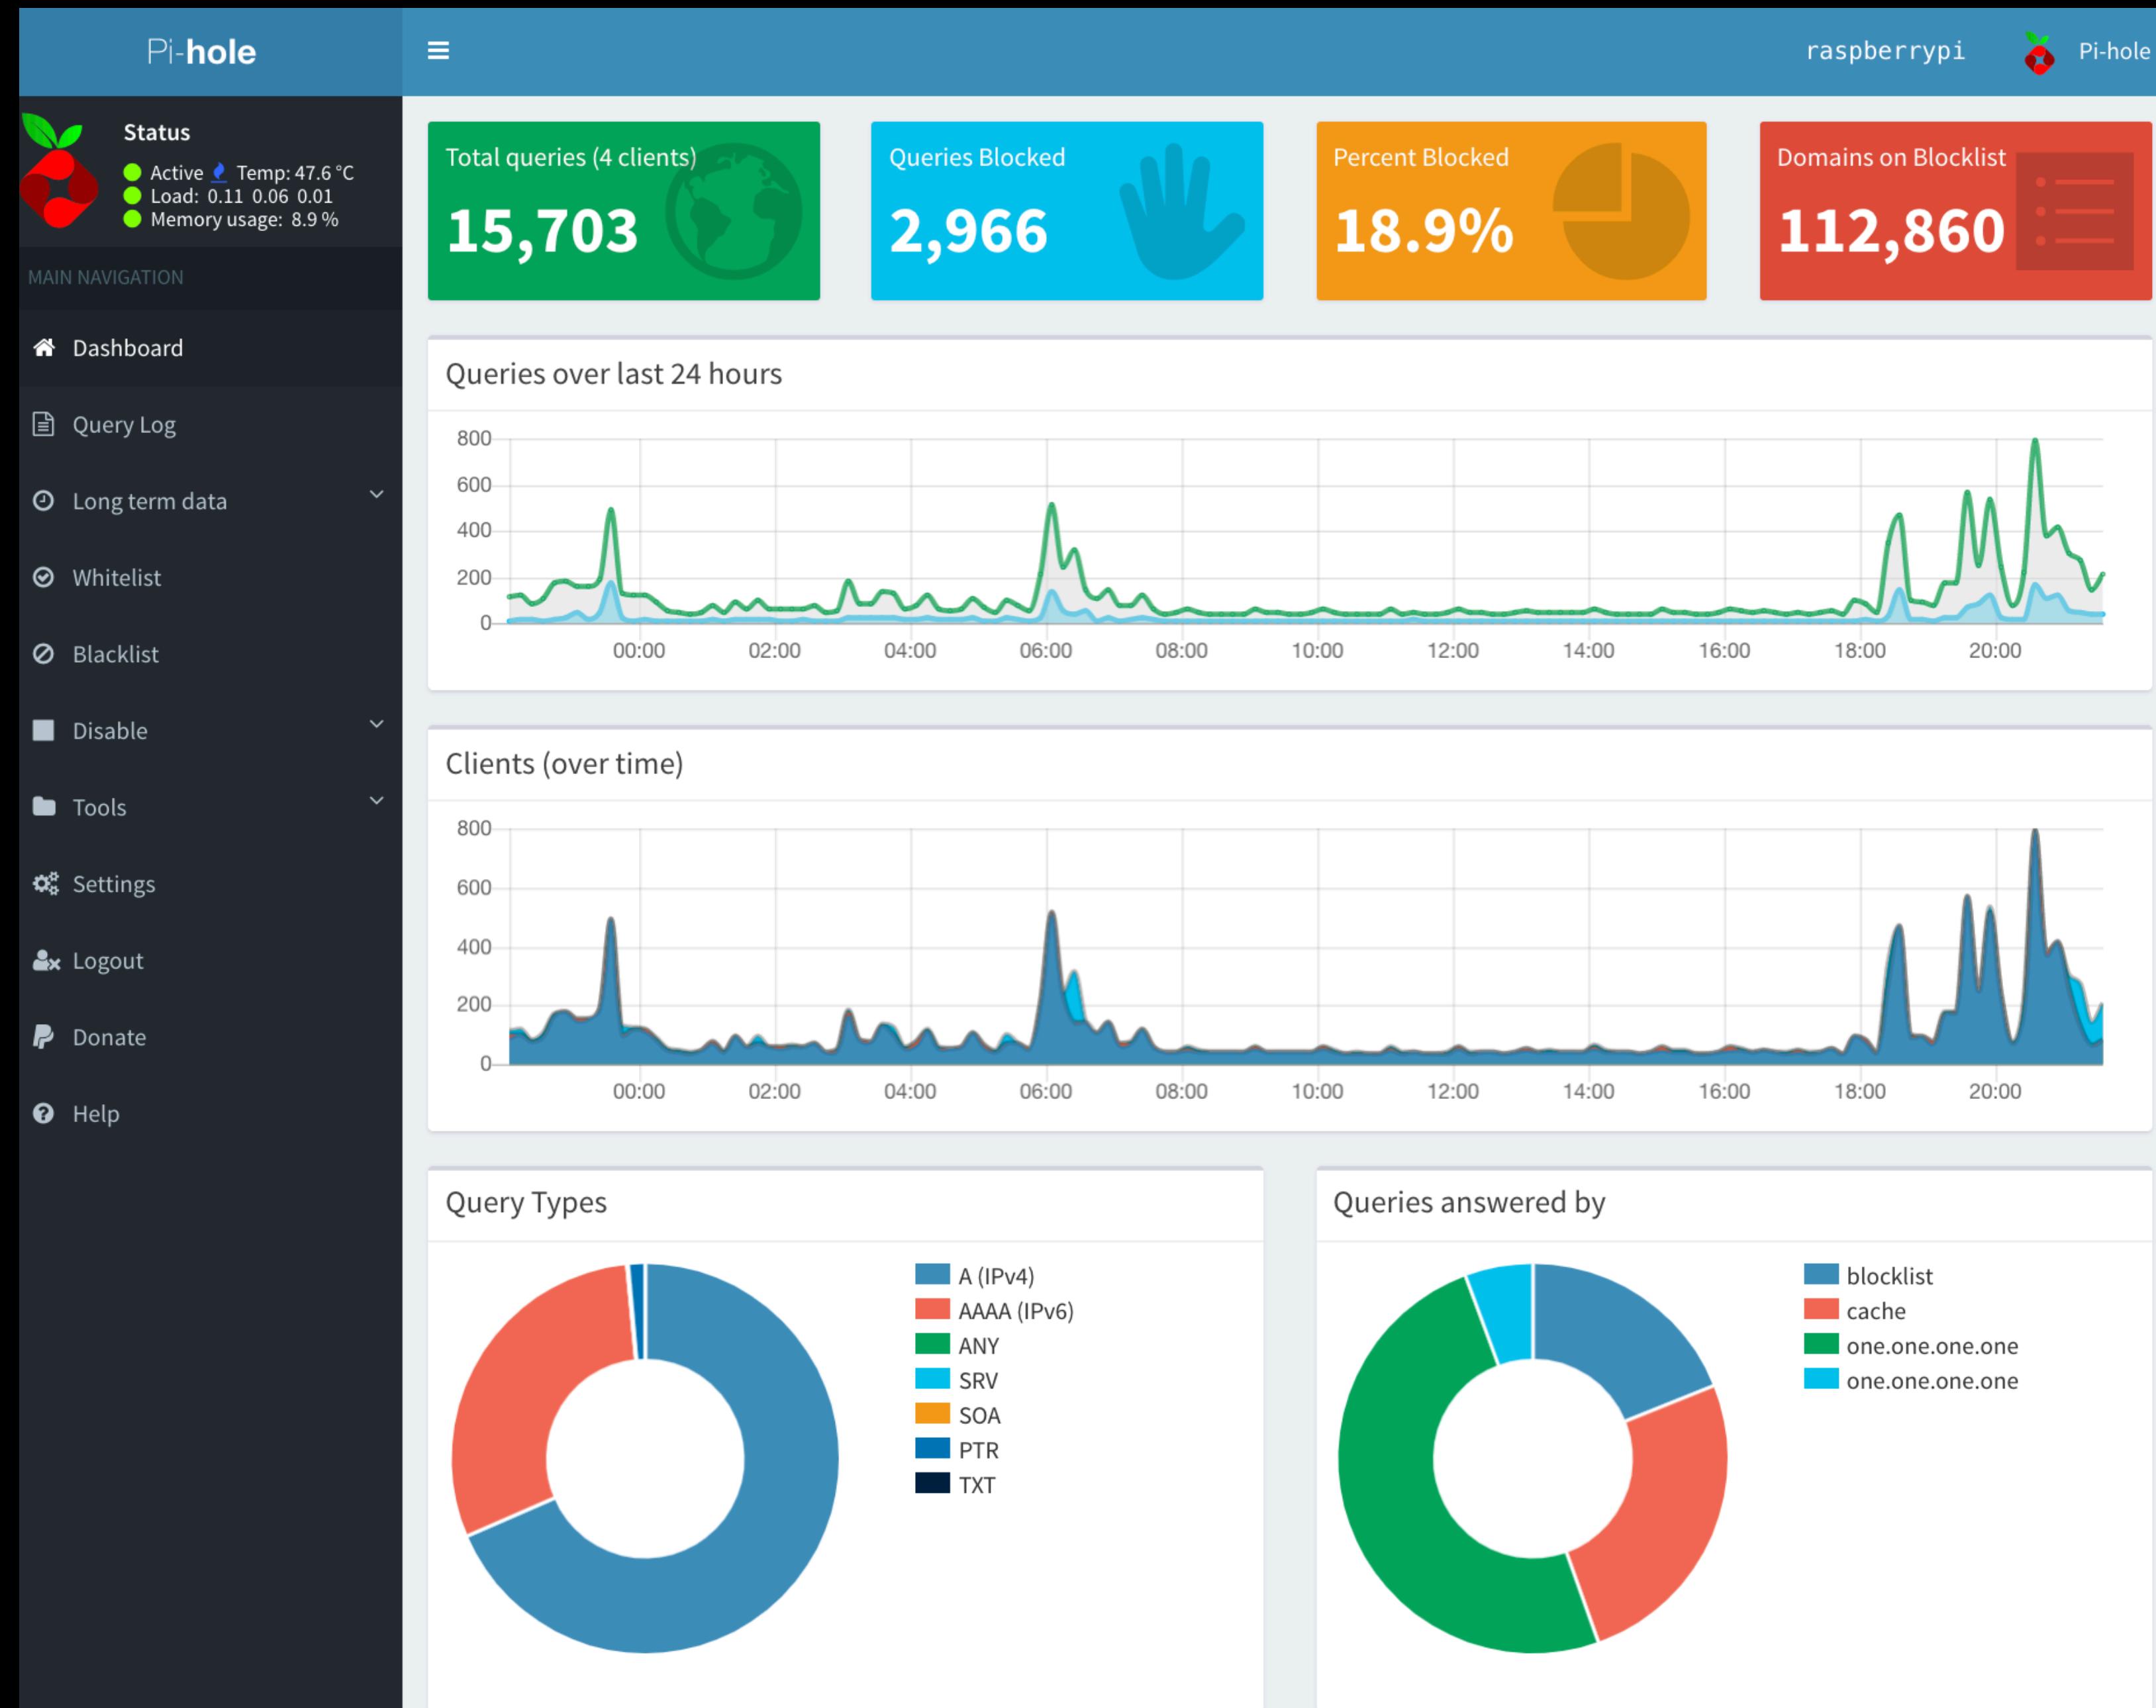


Monitor Everything/\*

20%



Don't Limit Monitoring to  
Infrastructure

- [ ] Introduction
- [ ] Operations Overview
- [ ] Failure Models
- [ ] Demo
- [ ] Best Practices & Recap

# docker run who-is-brian

## Brian Christner

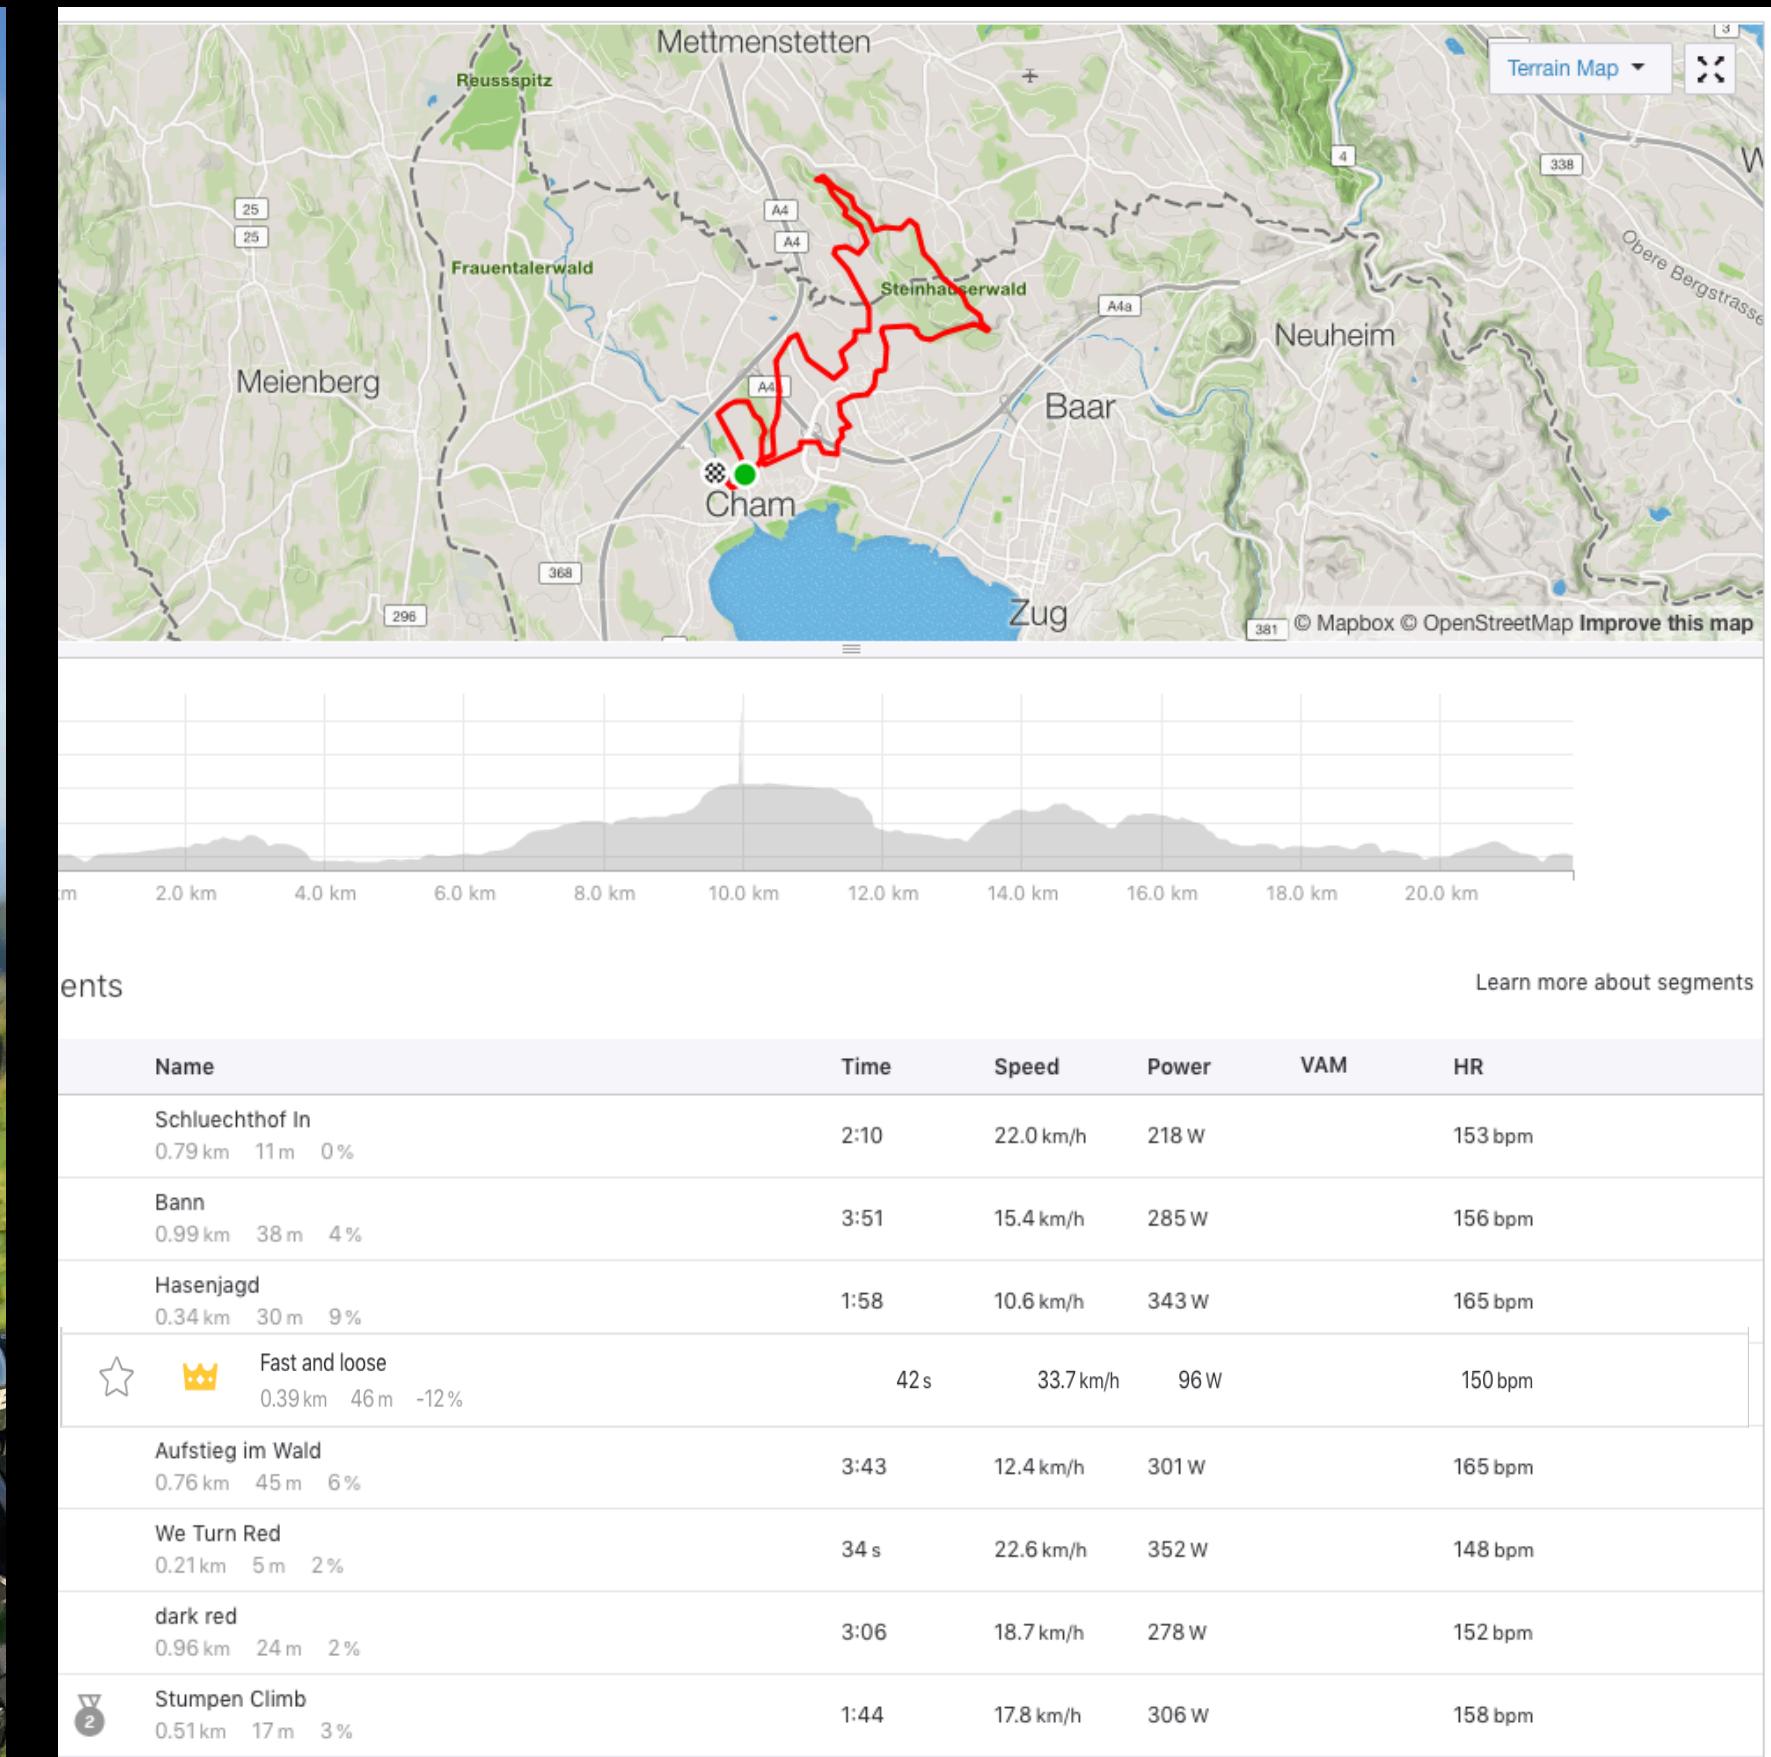
- Co-Founder 56K.Cloud / SRE
- Docker Captain
- Passionate about Monitoring and anything with .io domains







# Zugerberg, Zug, Switzerland





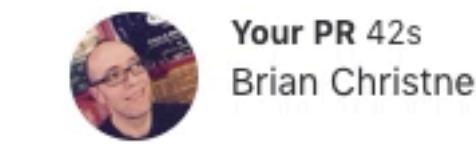
Fast and loose / Effort Comparison

## Effort Comparison

Share

## Fast and loose

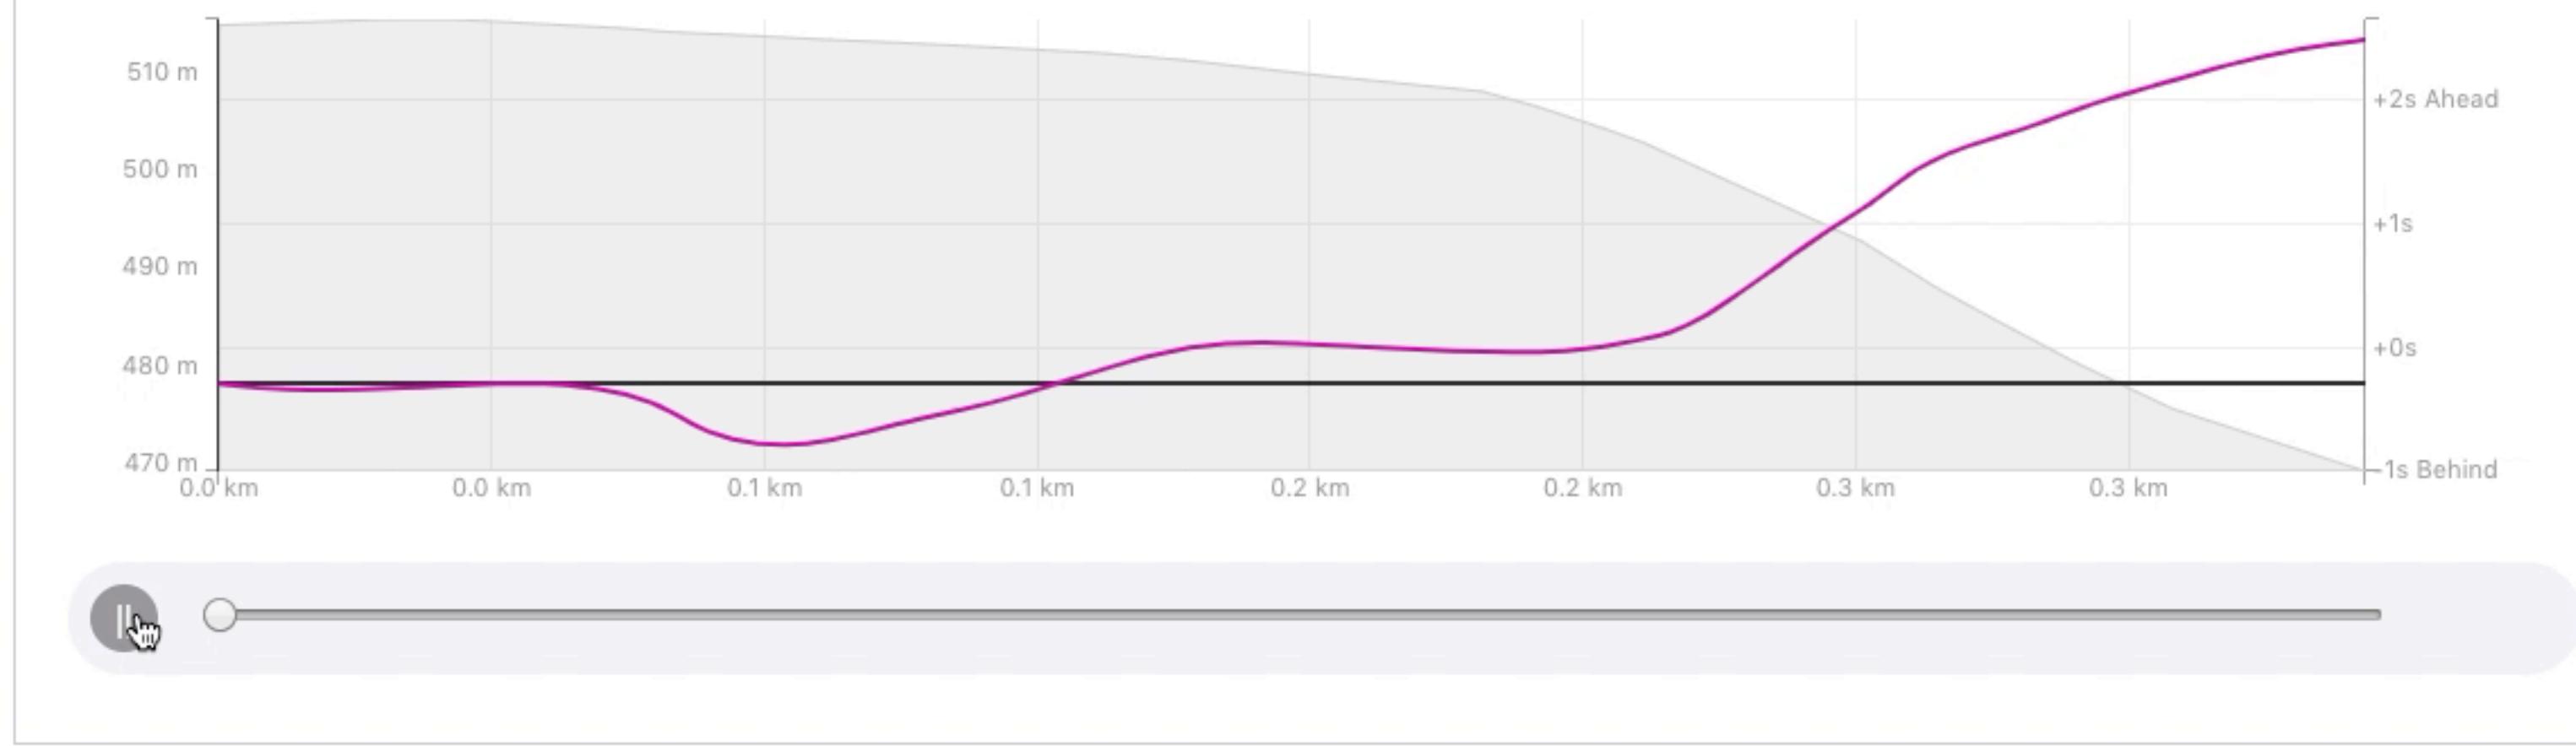
0.3km -12% Grade 46m

Your PR 42s  
Brian ChristnerKOM 39s  
Dani Mendlar

Athletes	Time	Speed
You 10/20/18 - PR	0s	36.0 km/h
Dani Mendlar 11/16/18 - KOM	0s	+4.0 km/h

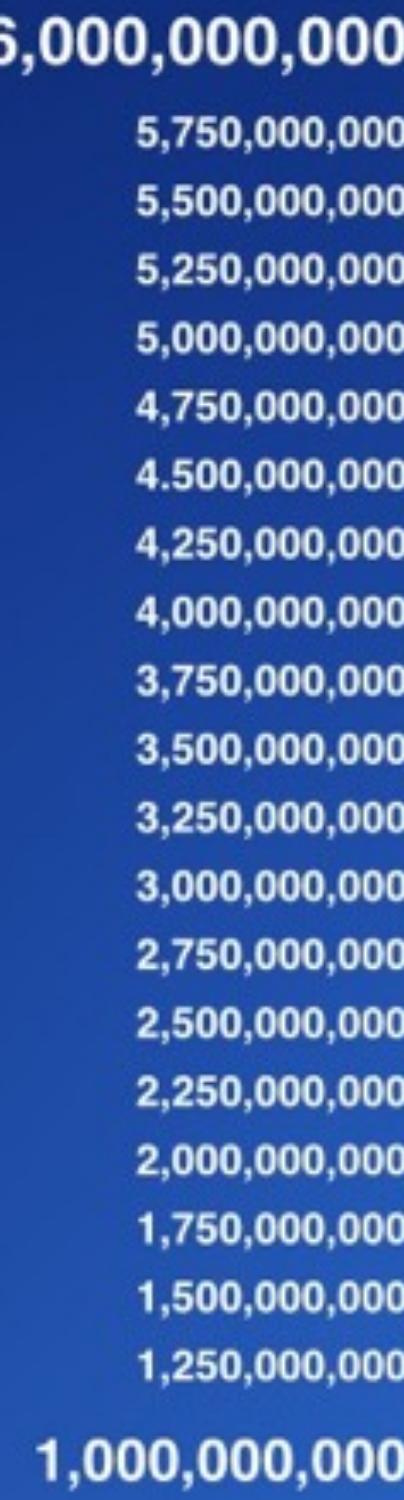
Customize and compare up to five efforts and see where the race was won.

[Get the Training Pack](#)



OK, give me a business  
example

# 2019 2 Billion++ every Week



HyperKit , VPNKit, DataKit

containerd  
Notary  
runC

libnetwork

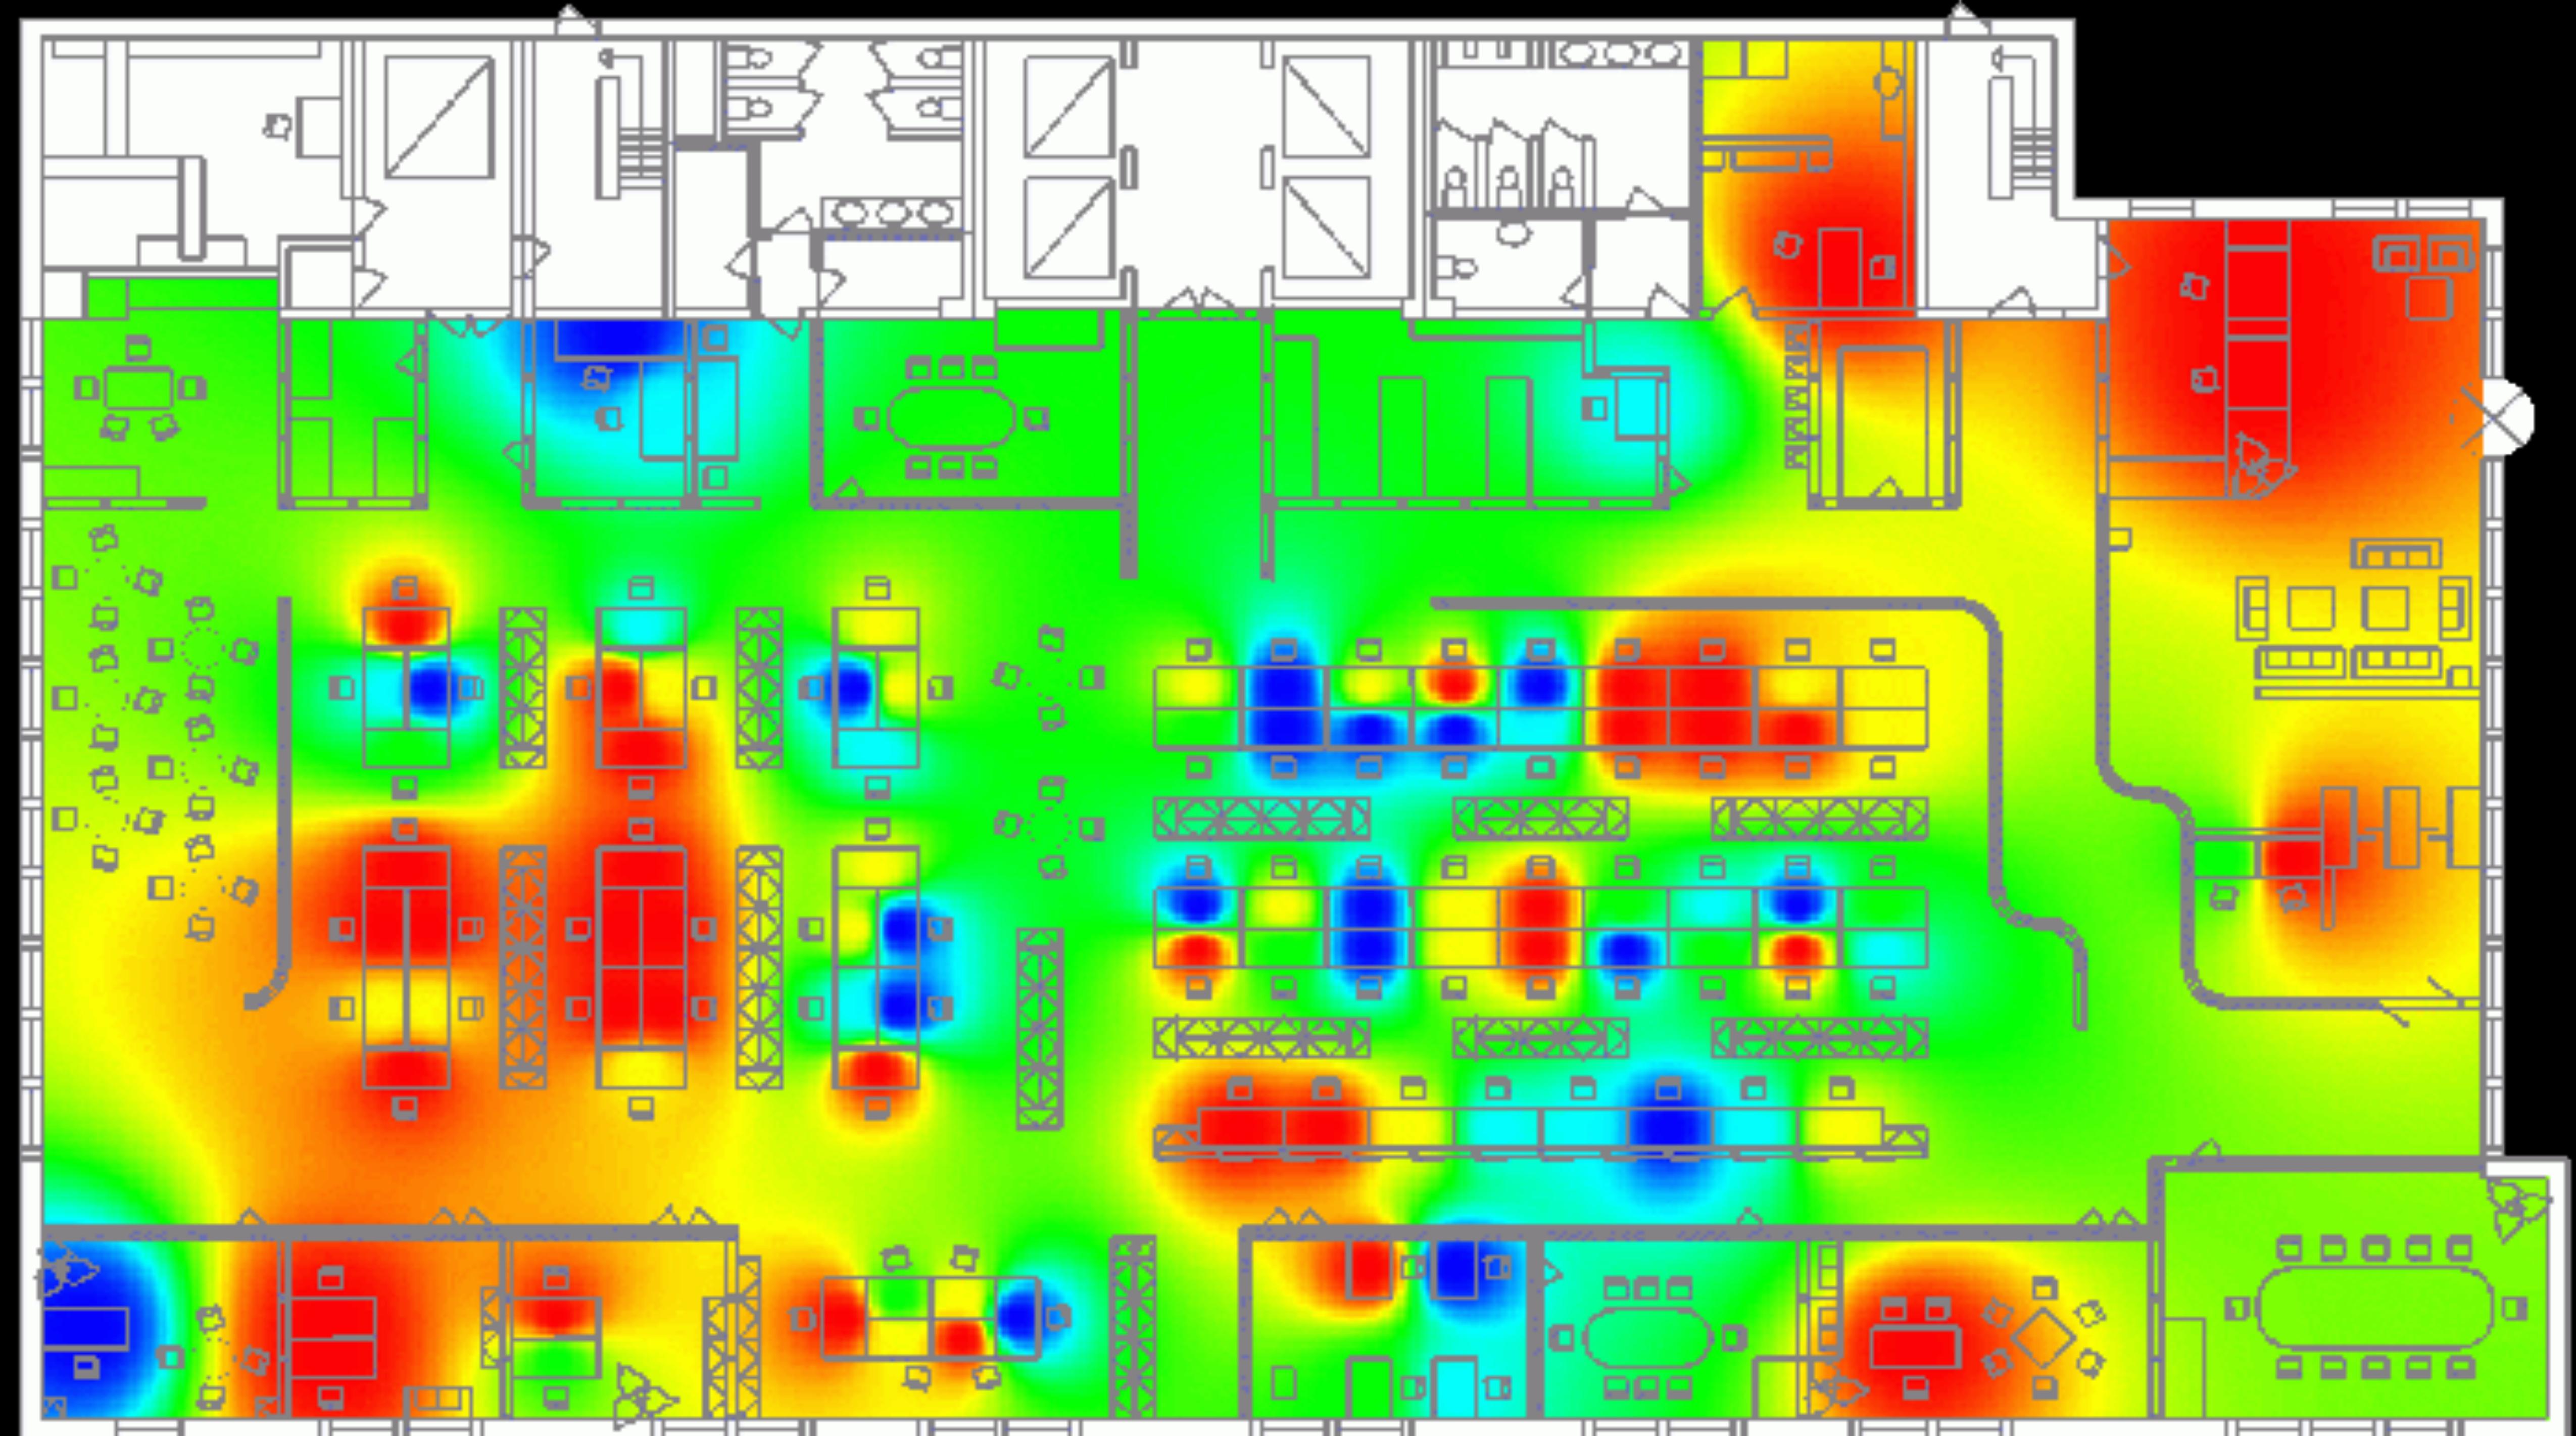
Docker 0.9 : Pluggable execution

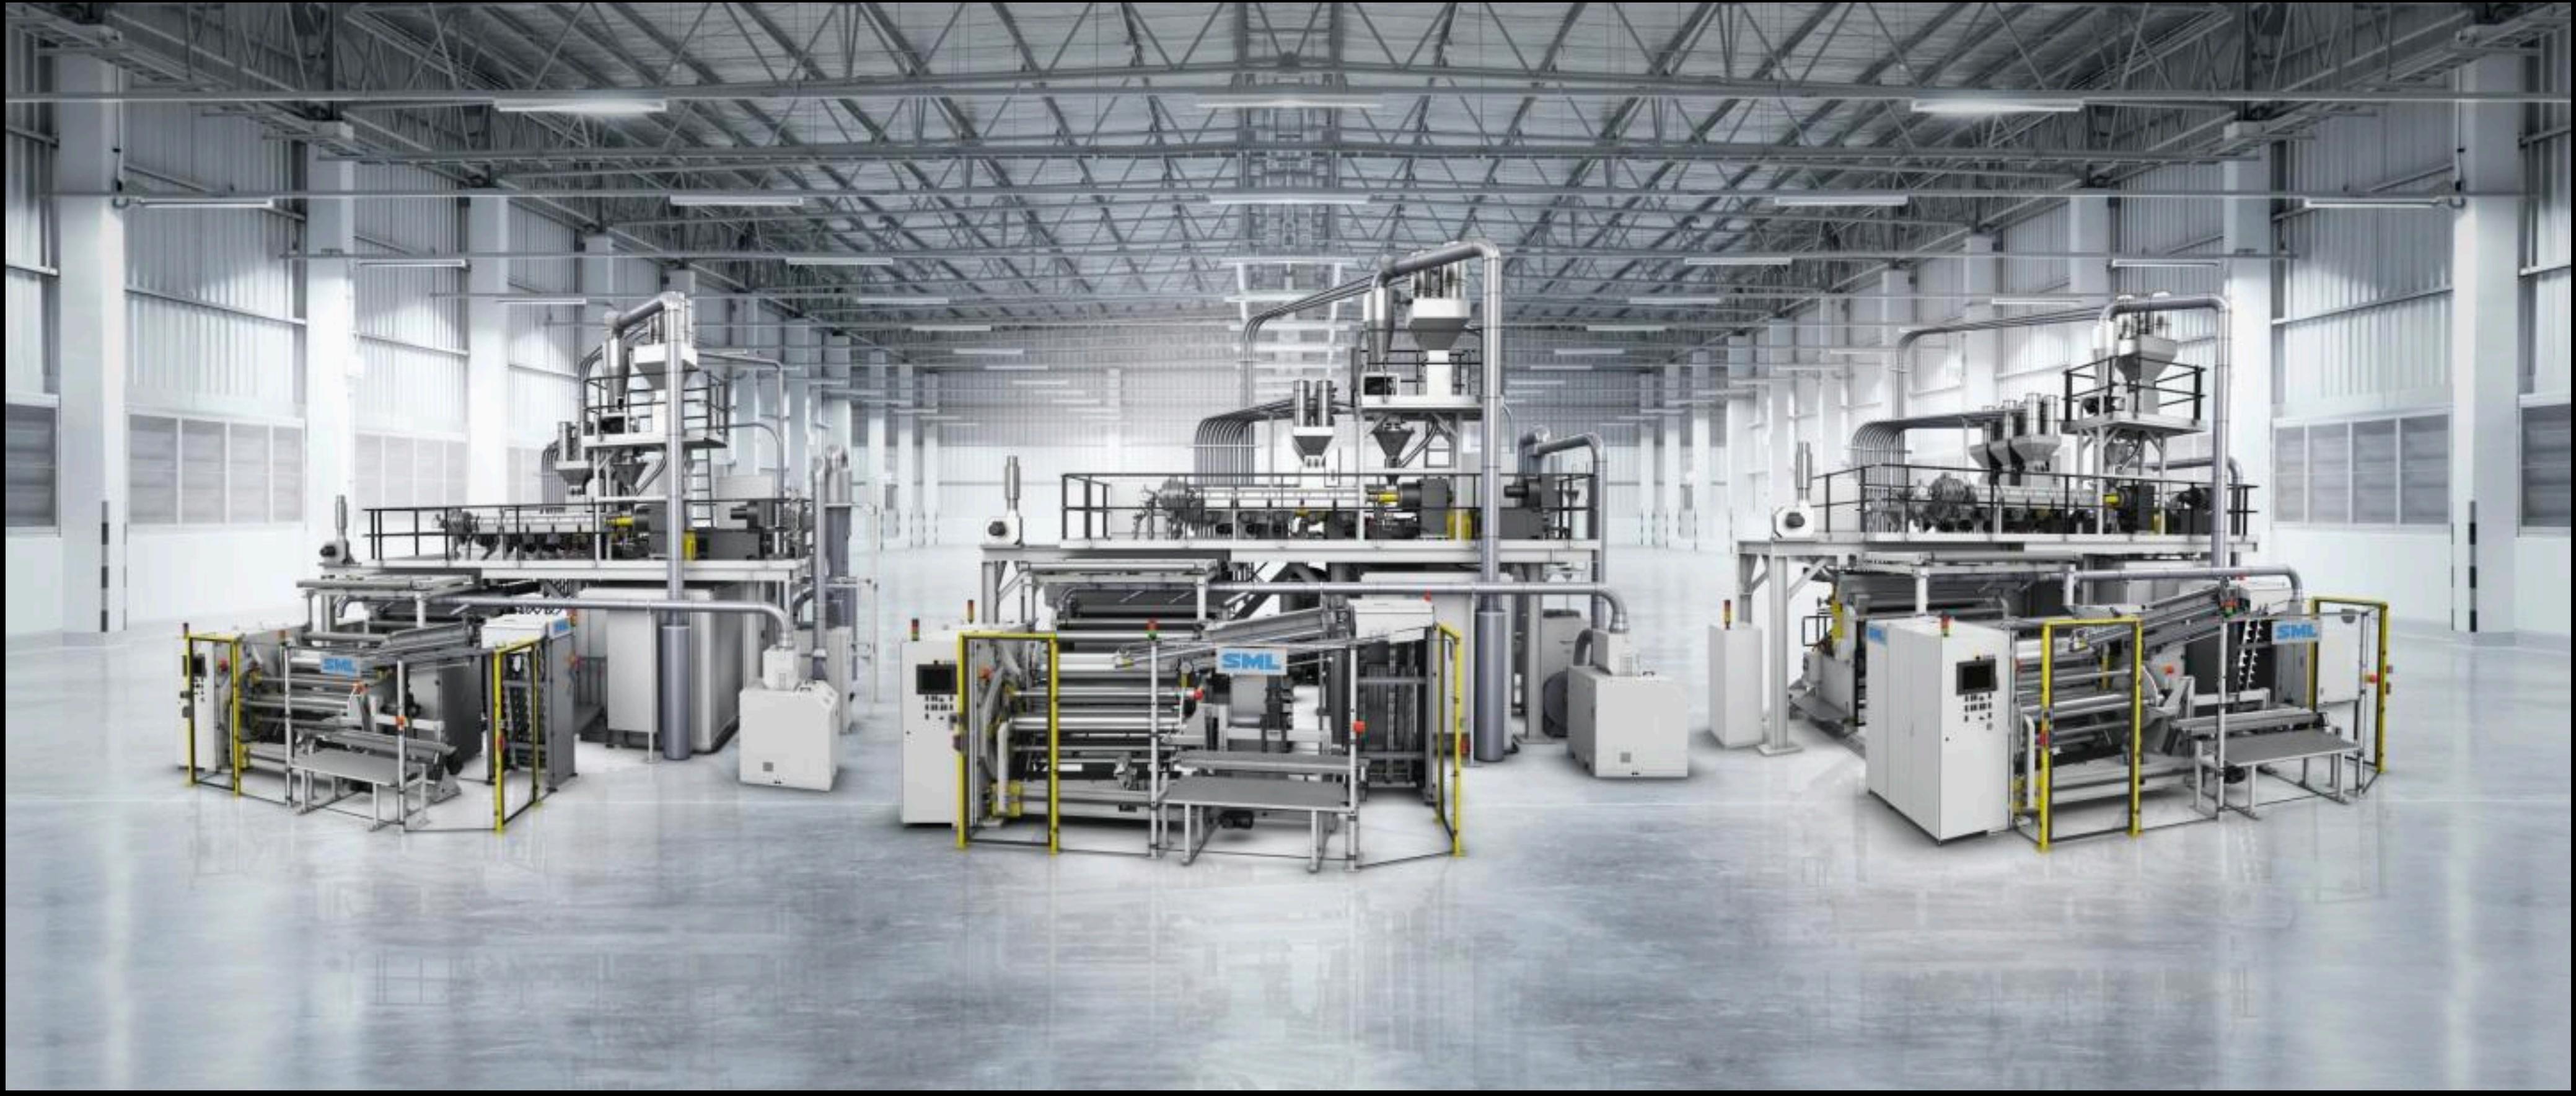
libcontainer

- SwarmKit
- Docker 1.12 with built-in orchestration
- Docker 1.11: OCI support
- Docker for Mac  
Docker for Windows
- Docker 1.8 : Docker Content Trust
- Docker 1.7 : Multi-Host Networking



# Casinos



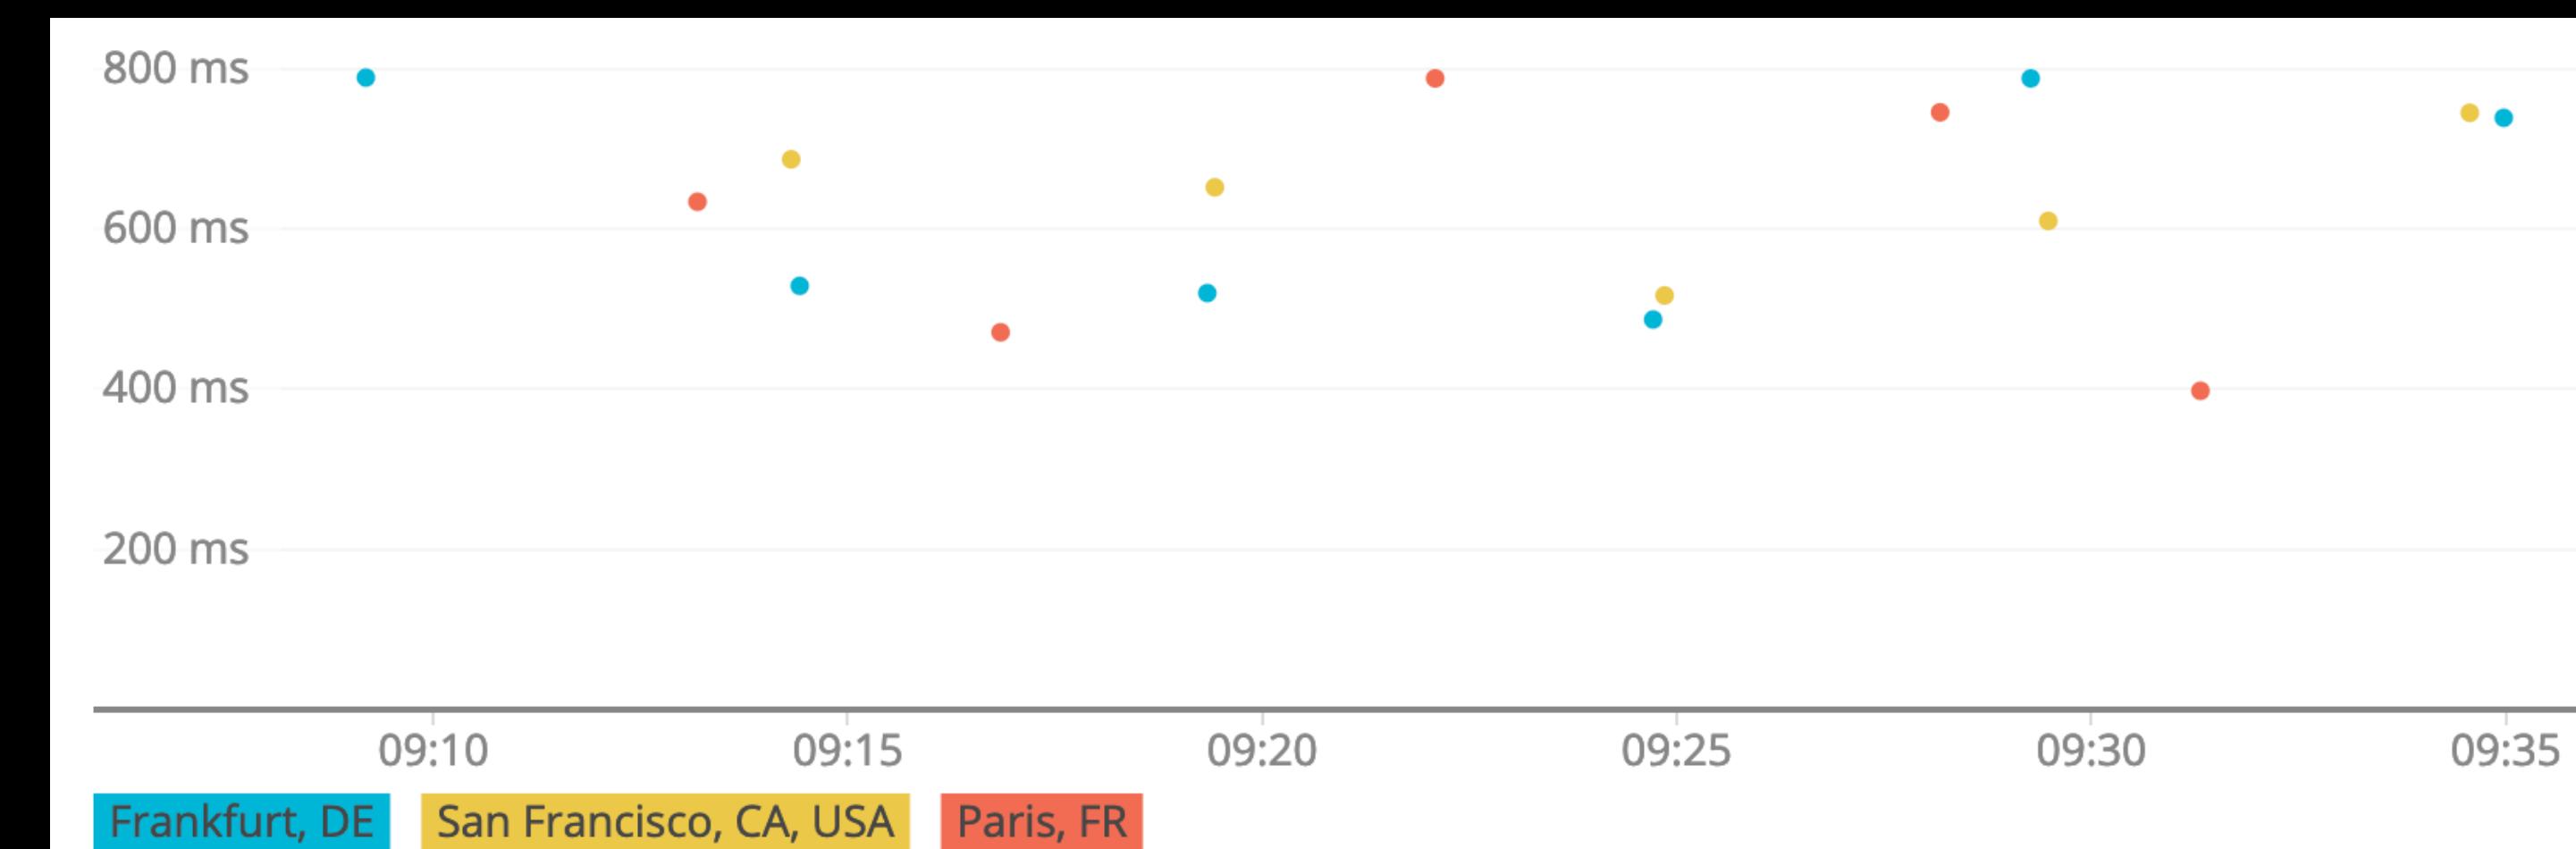
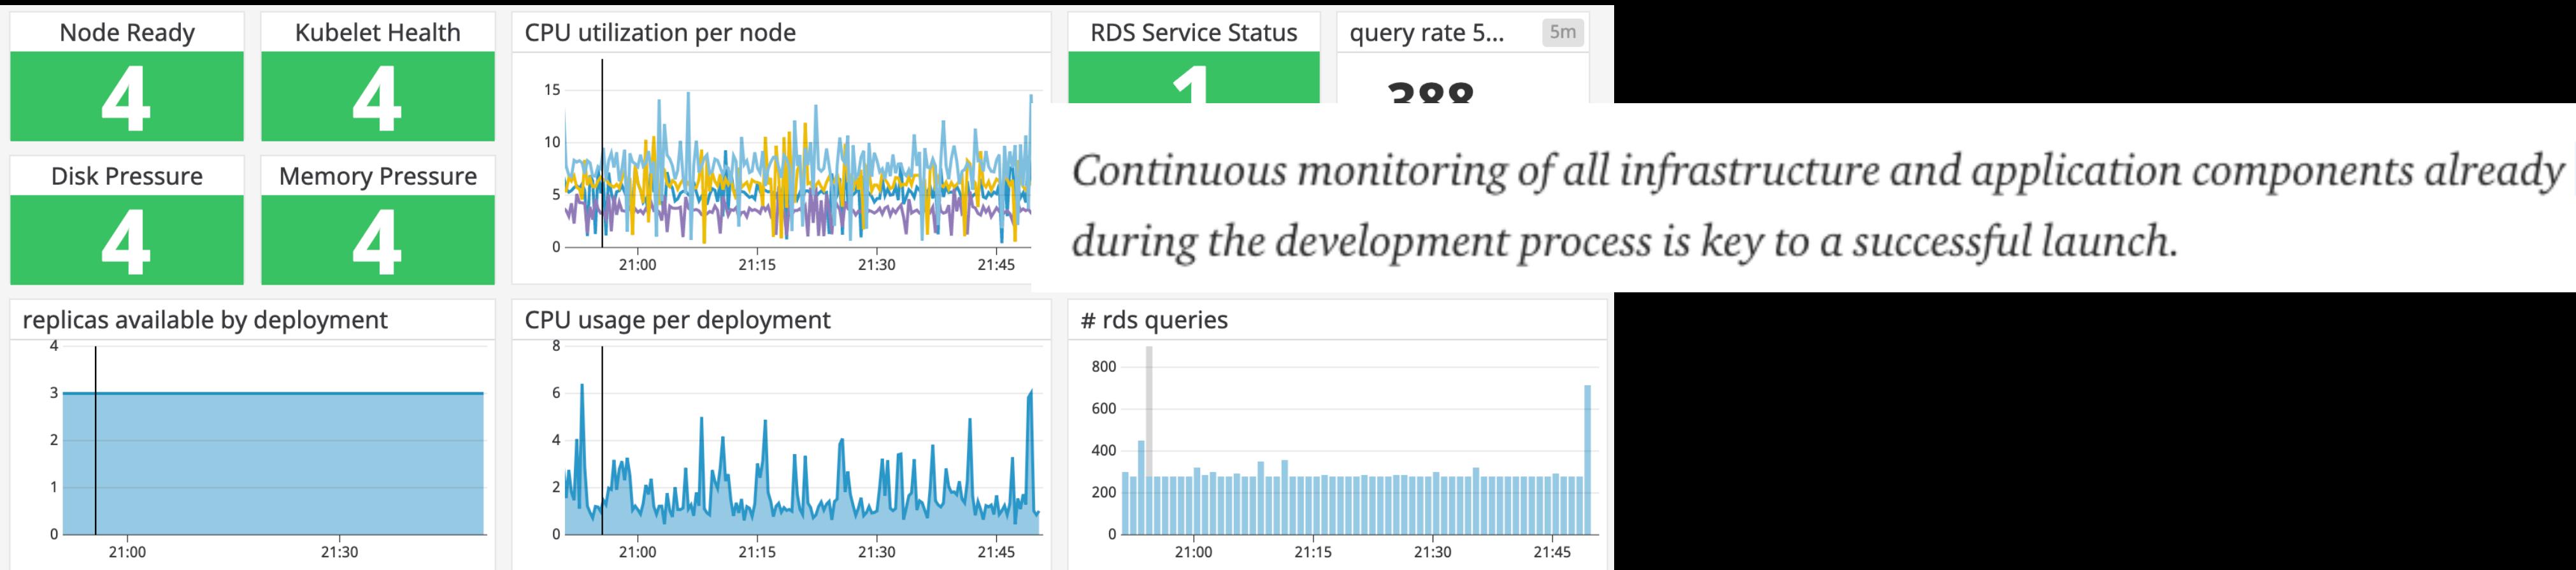


# Industrial machines

# Monitor Early

# www.your-now.com

A joint collaboration between BMW & Daimler AG



- [ X ] Introduction
- [ ] Operations Overview
- [ ] What to Monitor
- [ ] Demo
- [ ] Best Practices & Recap

# Ops Paradise



- Everything is Automated
- Reduce Costs
- No support calls / support tickets

# Ops Firefighting



# Operational Models

Manual	<ul style="list-style-type: none"><li>• User initiated</li><li>• Interactive, command-line tools, simple scripts</li><li>• Checklist and process driven</li></ul>
Reactive	<ul style="list-style-type: none"><li>• Hardware-centric data collection</li><li>• Simple metric and log collection</li><li>• Siloed tools and information</li><li>• Manual analysis and remediation</li></ul>
Proactive	<ul style="list-style-type: none"><li>• Application-centric data collection</li><li>• End-to-end <b>observability</b></li><li>• <b>Key metrics</b> and thresholds well understood</li><li>• Semi-automated analysis and remediation</li></ul>

# Users Care about 3 Things

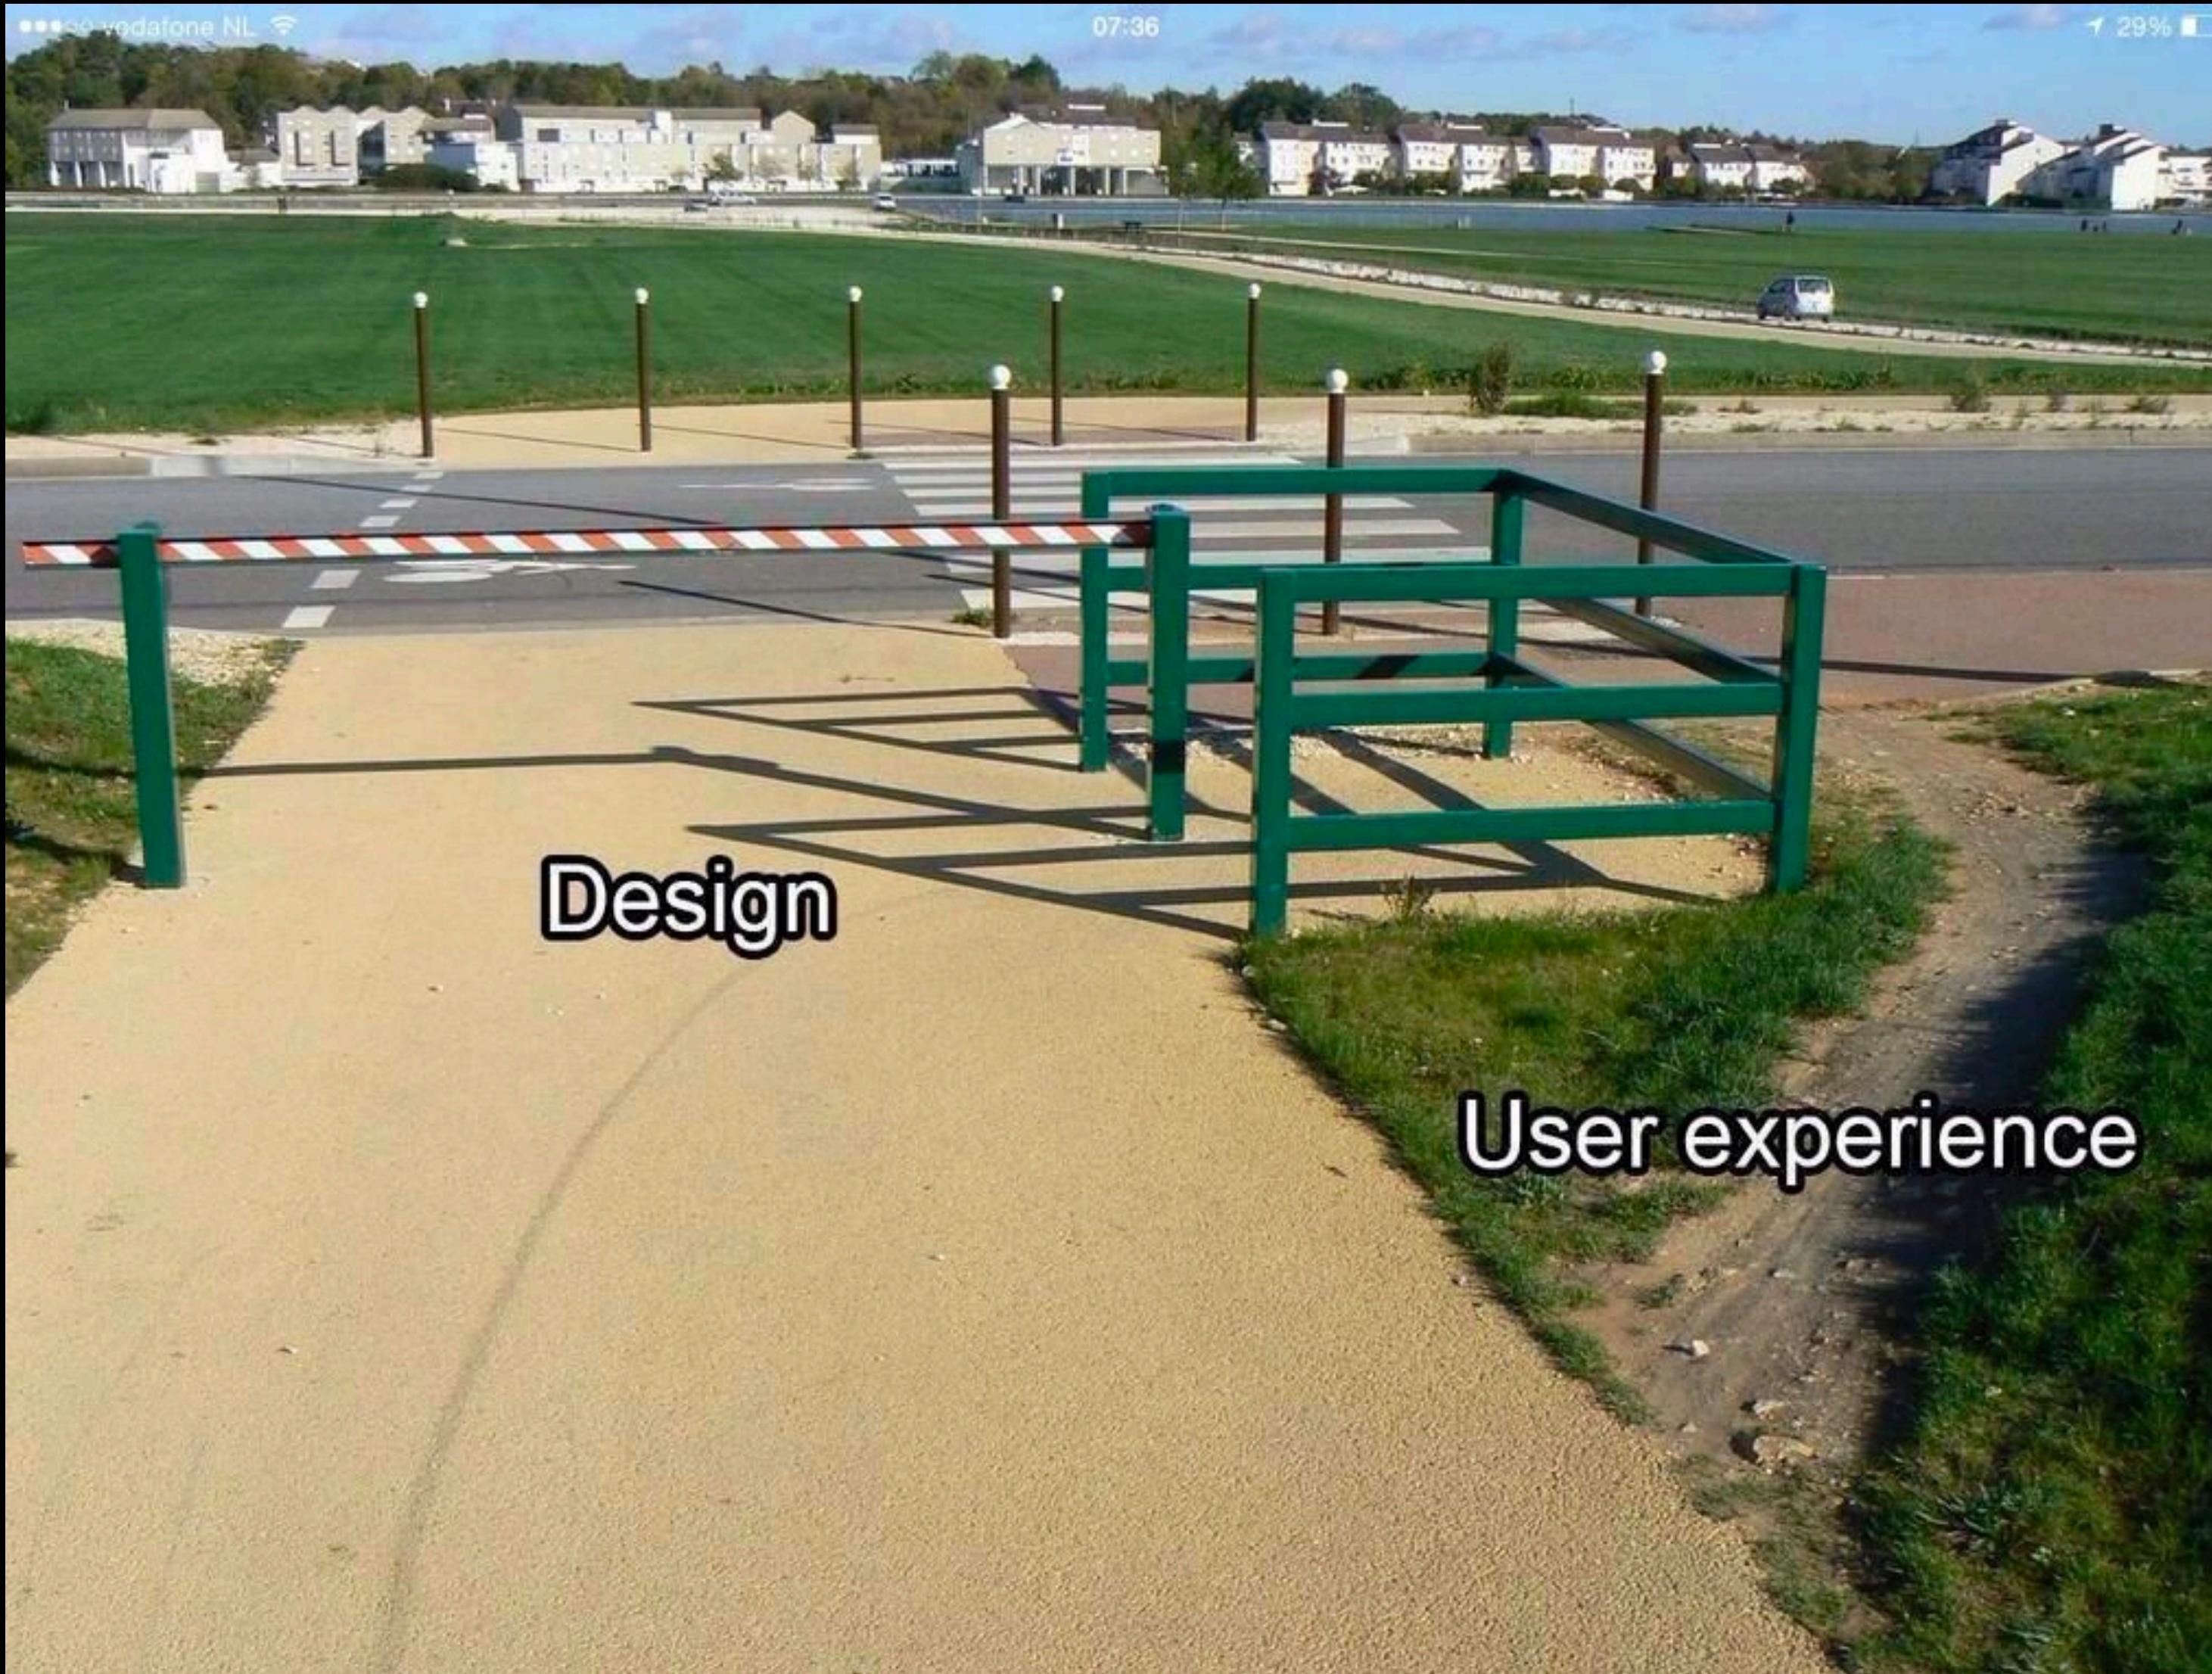
- **Availability** - Is my System Online Yes/No
- **Latency** - Does it take a long time to access applications x, y, z
- **Reliability** - Can the user rely on using the application

# Brain Based Tools



- We can track 8 objects on average
- 4 Moving Objects
- Build Dashboards & Tools accordingly

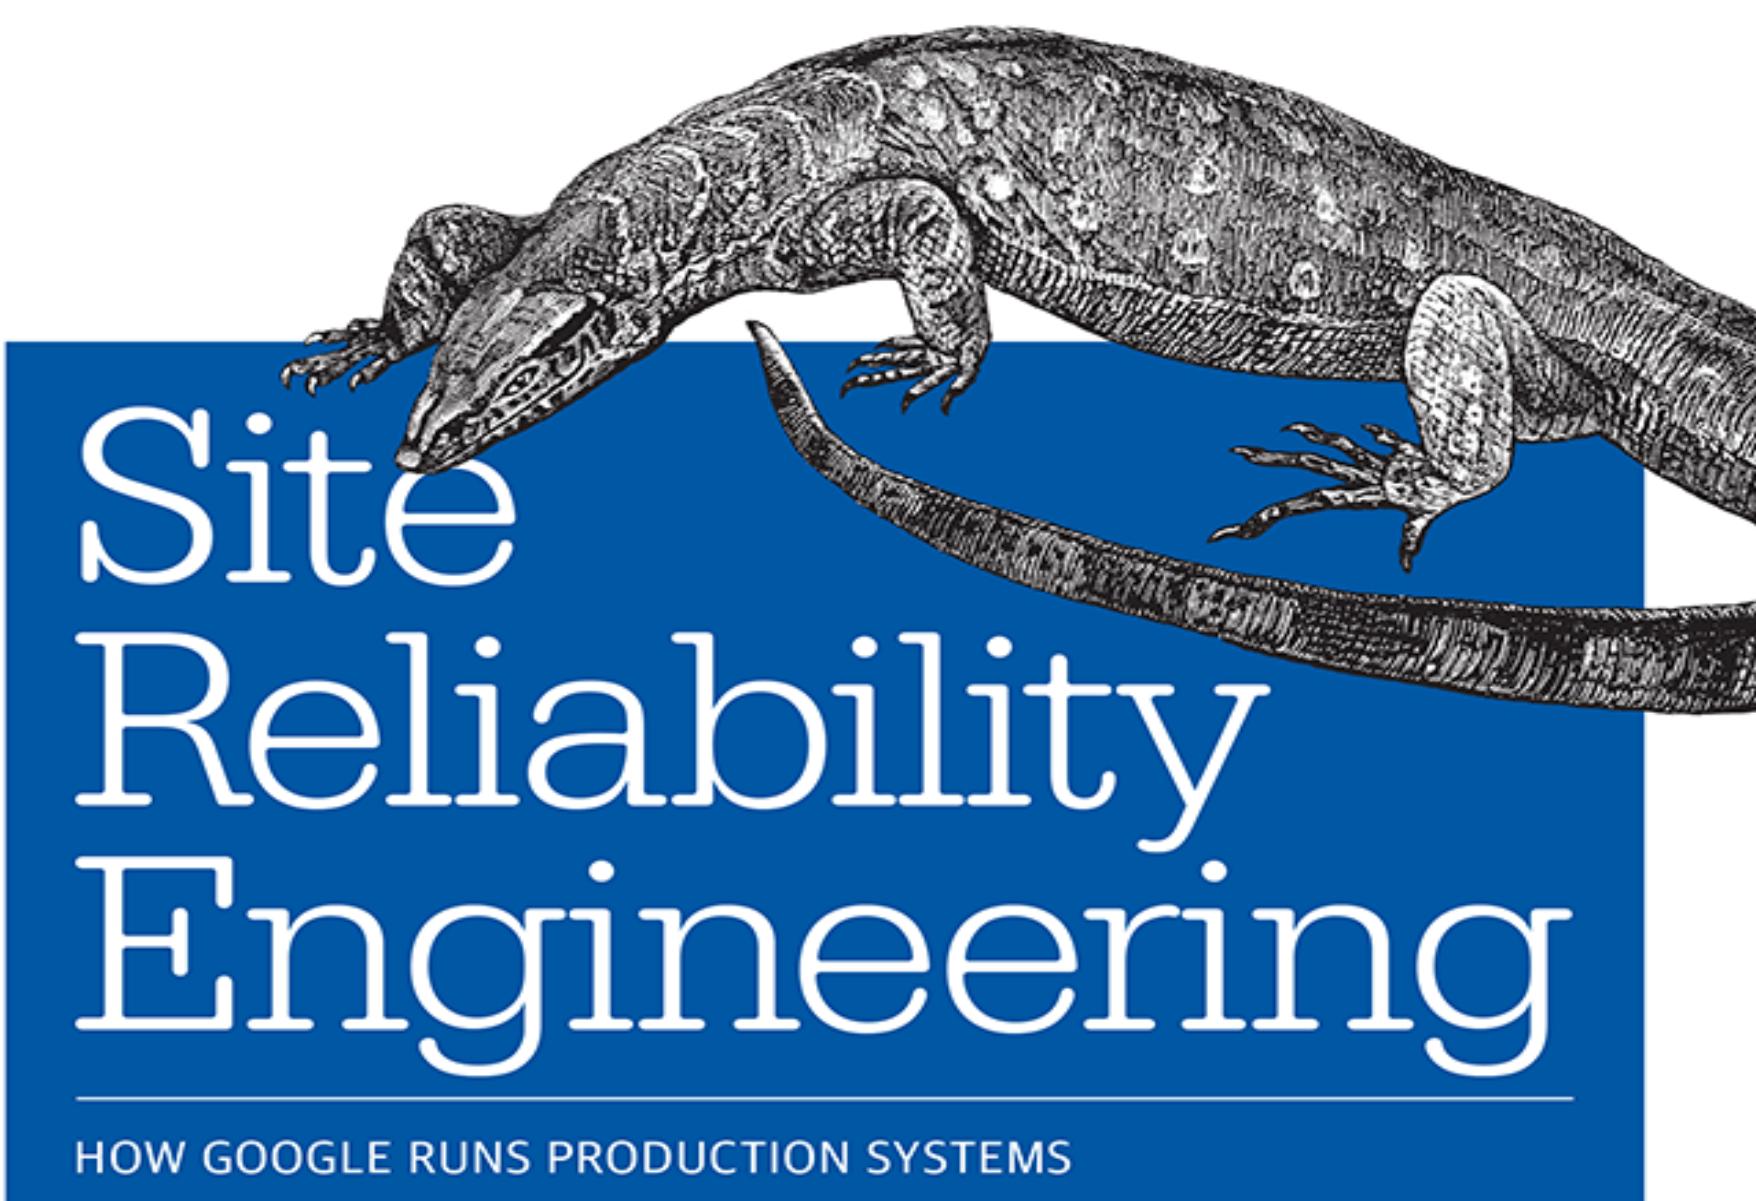
# Design vs Experience



# SRE

# SRE (Site Reliability Engineering)

O'REILLY®



SRE is treats Operations as if it  
were a Software Problem

*“Hope is not a strategy.”*  
Traditional SRE saying

Edited by Betsy Beyer, Chris Jones,  
Jennifer Petoff & Niall Murphy

[www.google.com/sre](http://www.google.com/sre)

# TL;DR - SRE

O'REILLY®



Site  
Reliability  
Engineering

HOW GOOGLE RUNS PRODUCTION SYSTEMS

Edited by Betsy Beyer, Chris Jones,  
Jennifer Petoff & Niall Murphy

## 4 Golden Signals



Latency

Traffic

Errors

Saturation

# R.E.D (Microservice Level)

---

**(Request) Rate:** the number of requests, per second, your services are serving.

---

**(Request) Errors:** the number of failed requests per second. Utilization: the average time that the resource was busy servicing work

---

**(Request) Duration:** distributions of the amount of time each request takes.

# U.S.E (Low Level / Infrastructure)

For every resource, check Utilization, Saturation, and Errors

---

**Resource:** all physical server functional components (CPUs, disks, busses, ...)

---

**Utilization:** the average time that the resource was busy servicing work

---

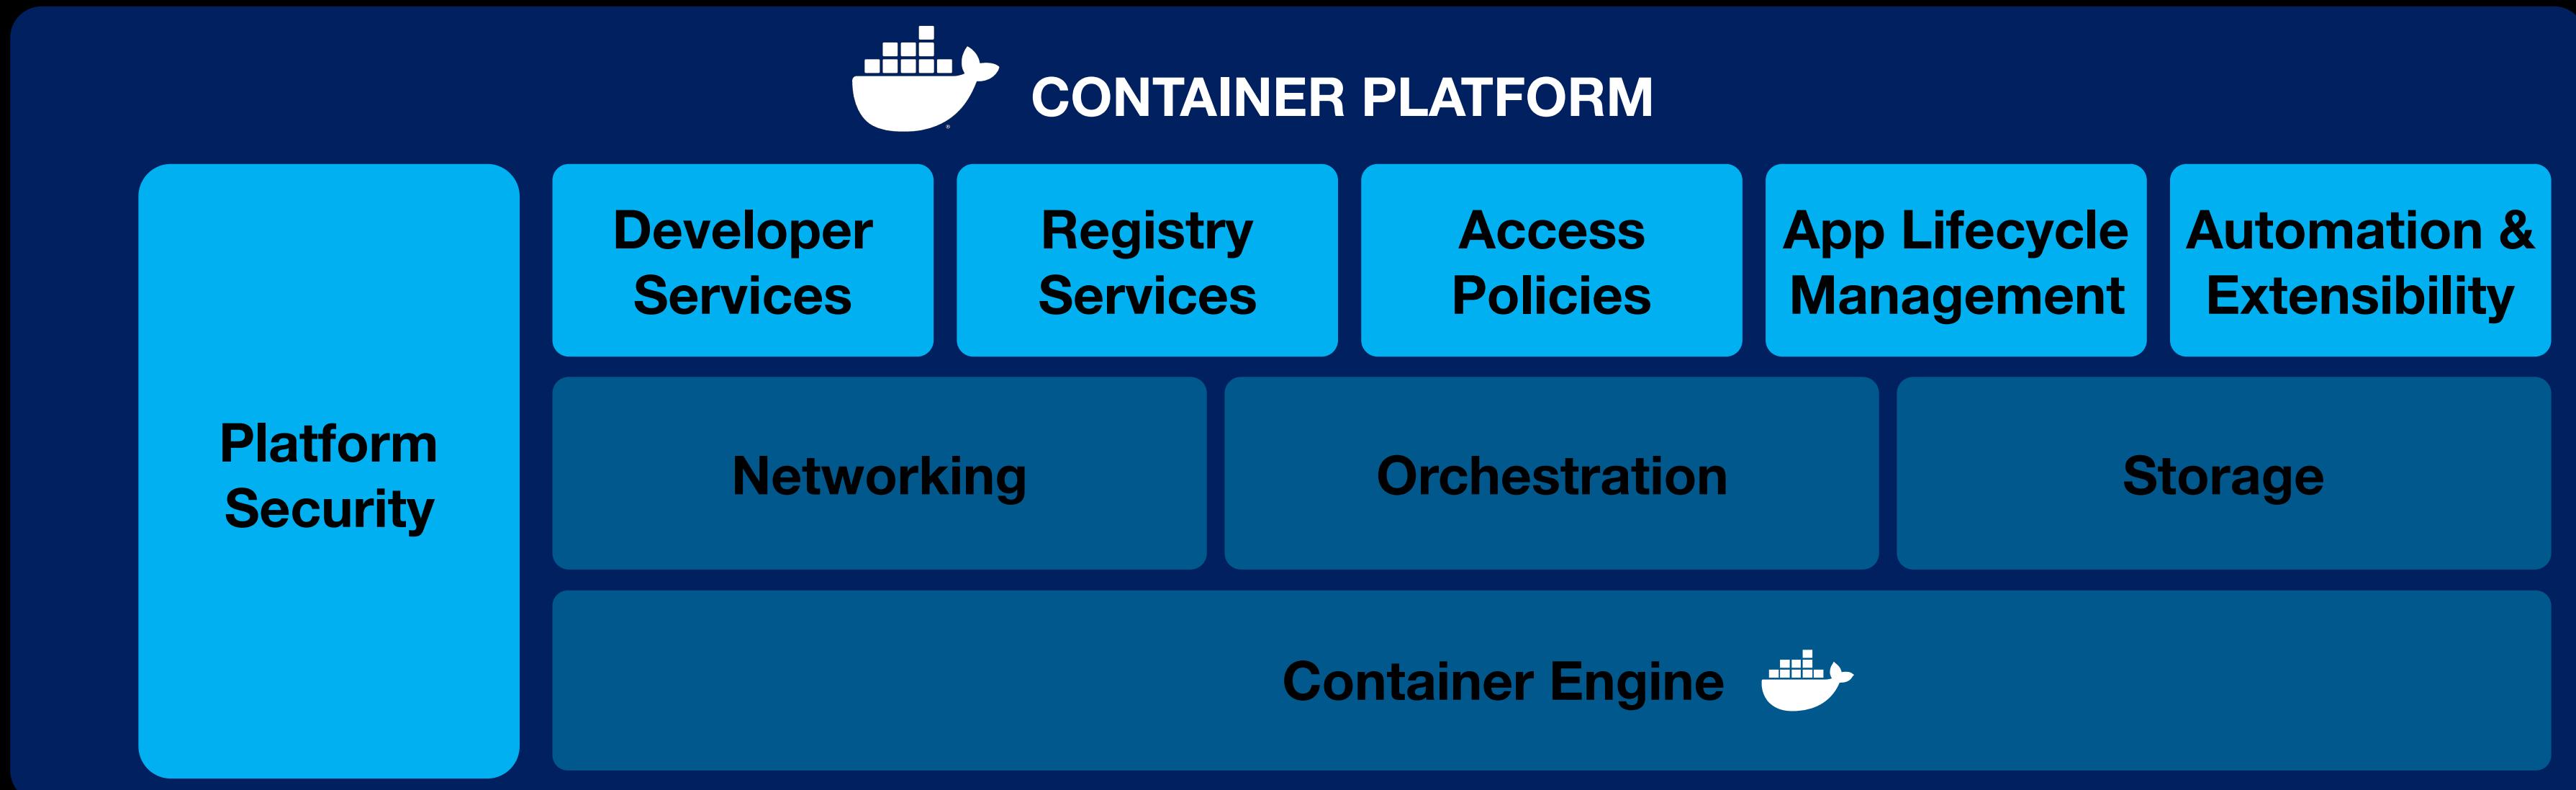
**Saturation:** the degree to which the resource has extra work which it can't service, often queued

---

**Errors:** the count of error events

- [ X ] Introduction
- [ X ] Operations Overview
- [ ] What to Monitor
- [ ] Demo
- [ ] Best Practices & Recap

# Understanding Failure Models



# Host / Hardware



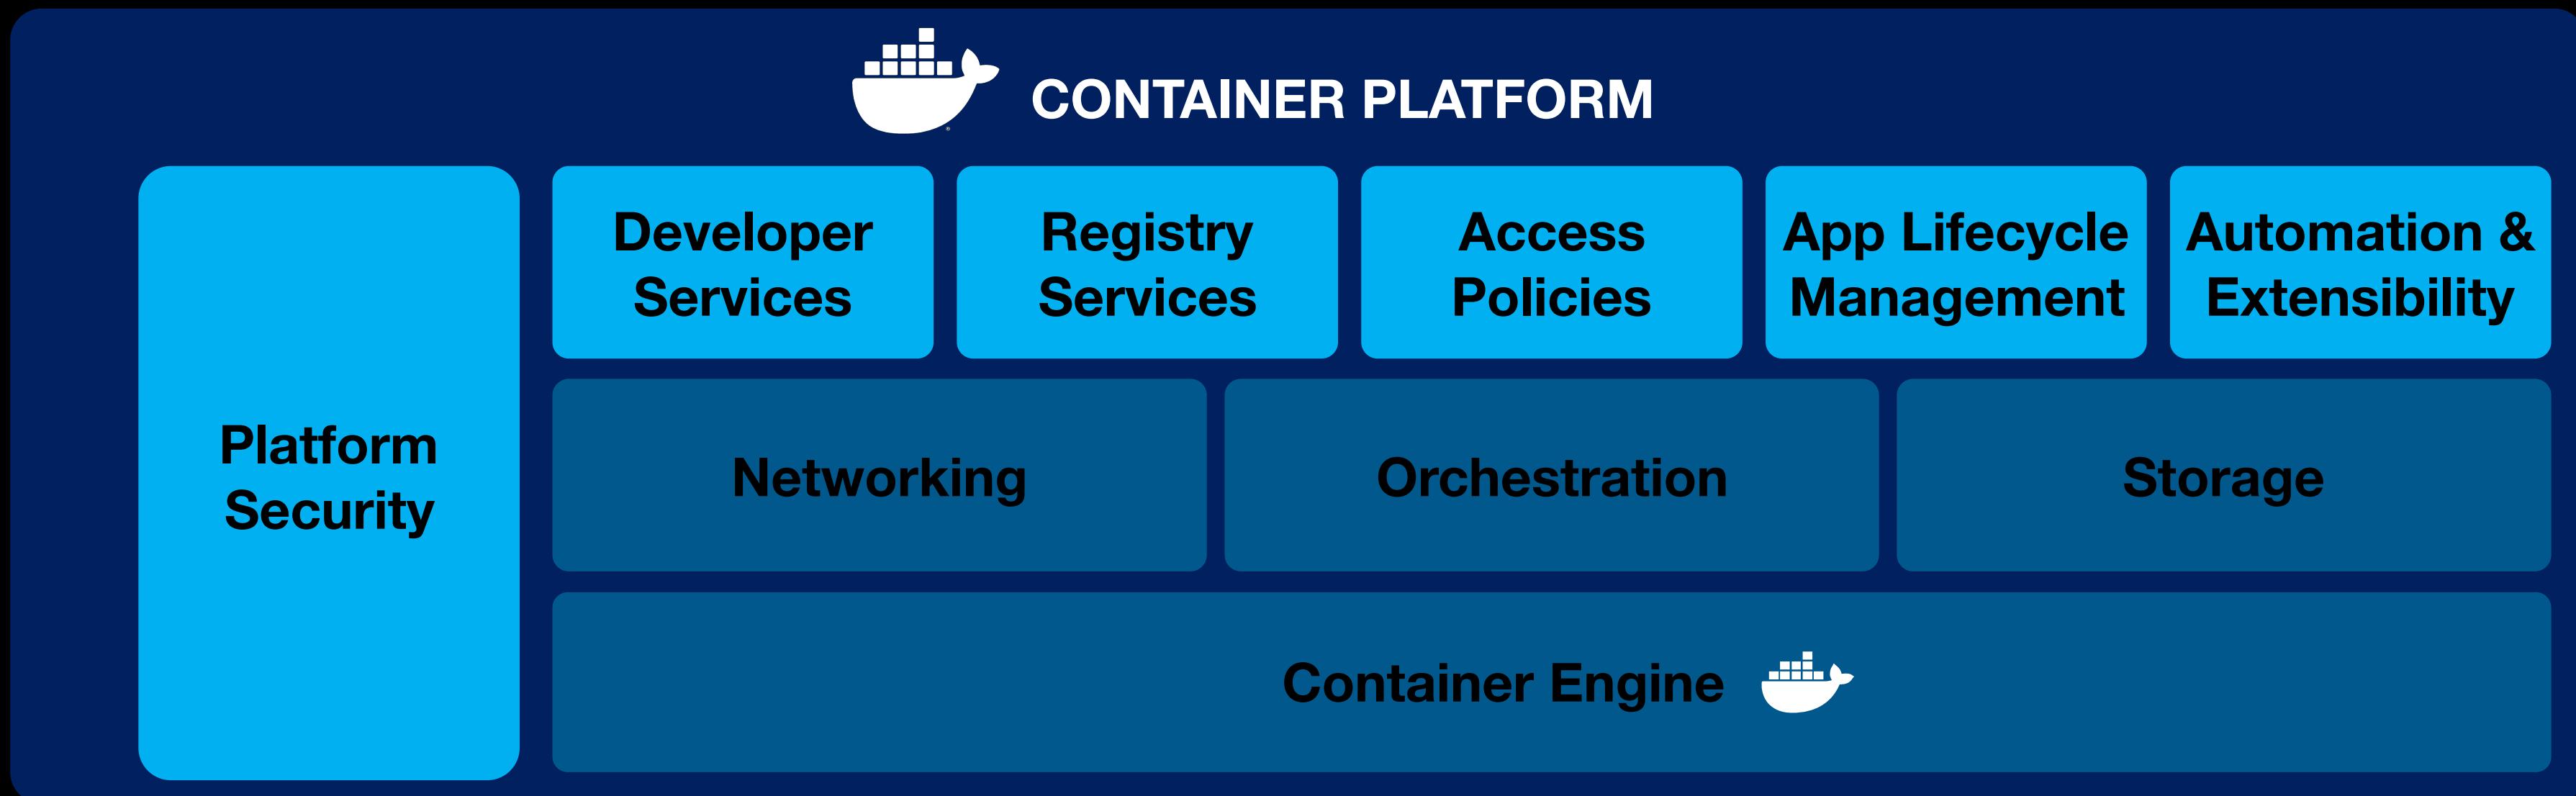
CPU

Memory

Liveness

File Descriptors

Storage Capacity



# Networking

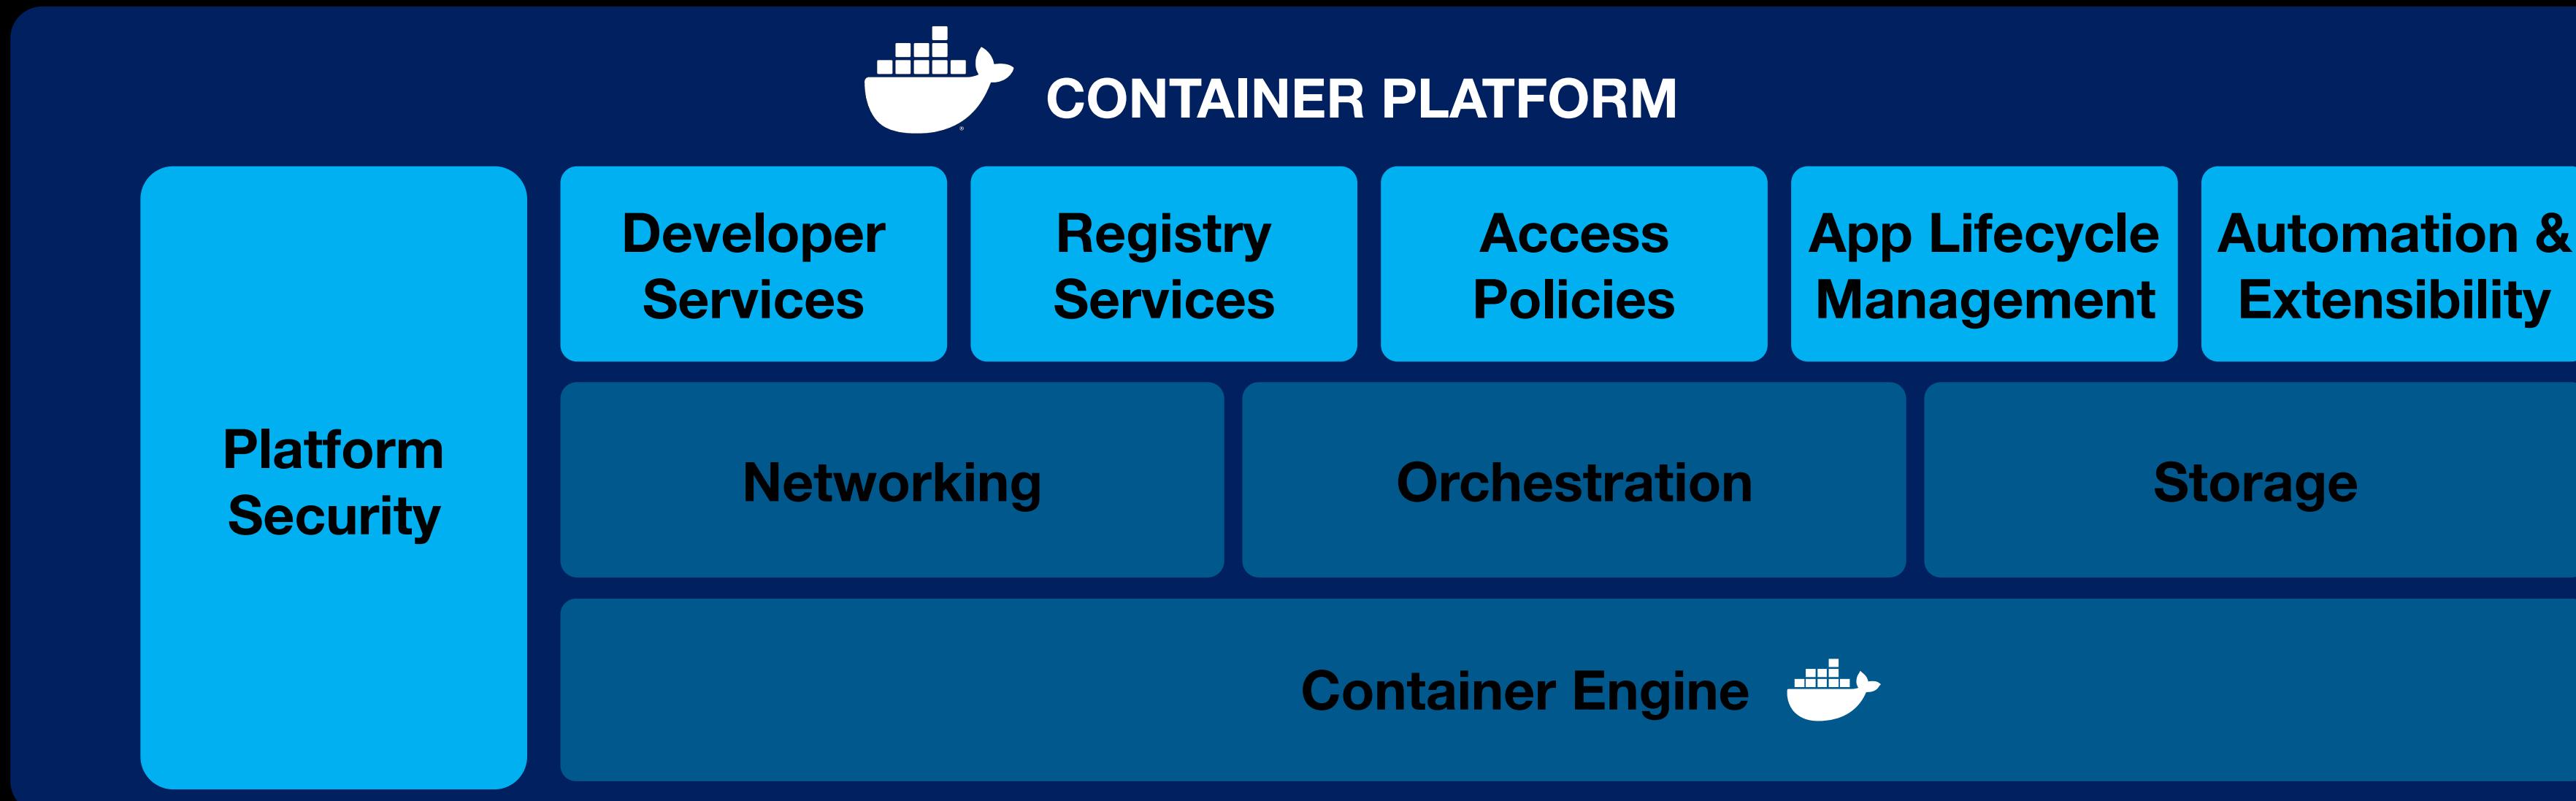


Reachability

Link Utilization

File Descriptors

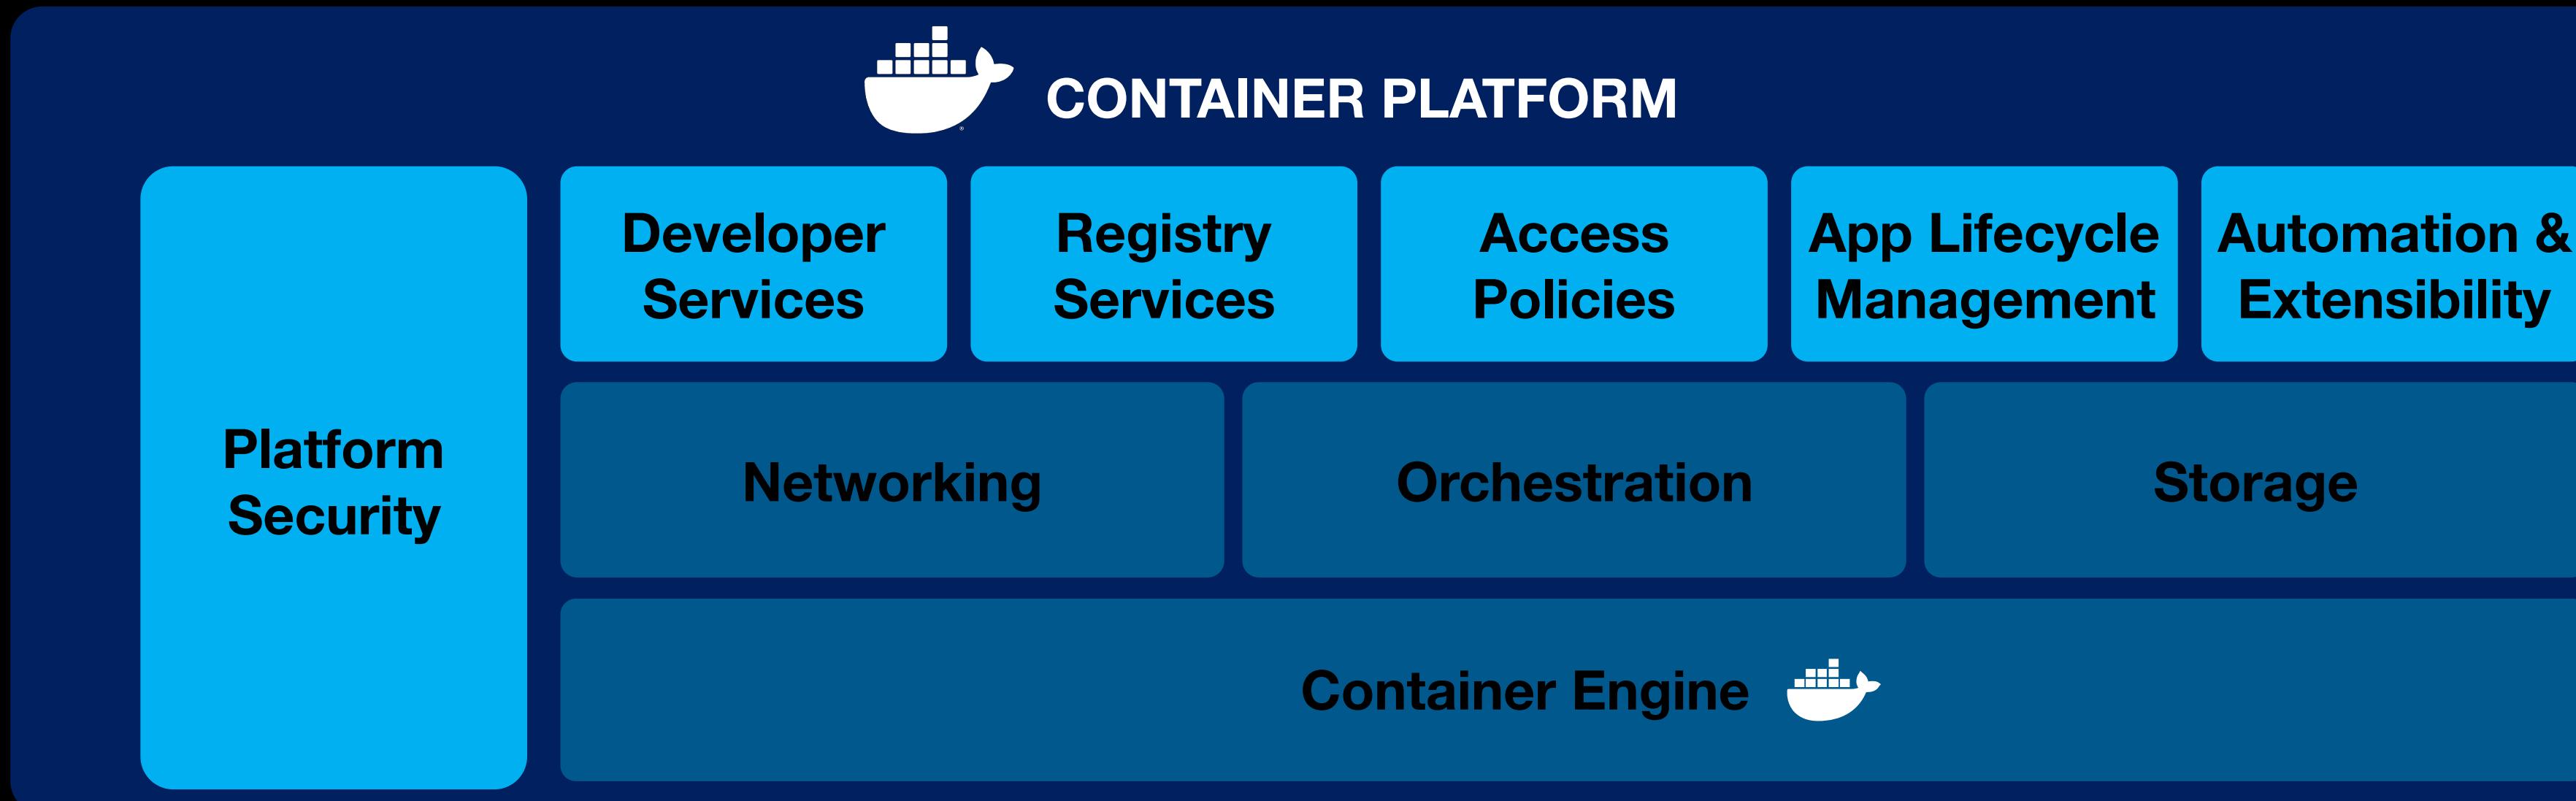
Storage Capacity



# Orchestration



State  
Deployment Rates  
Capacity  
Scheduling Events



# Applications



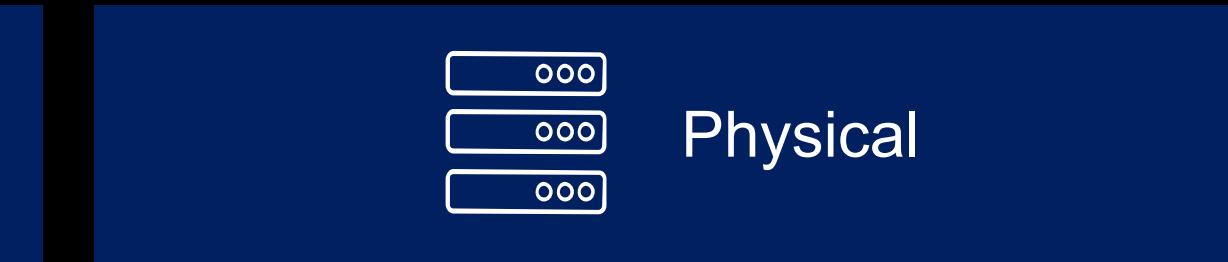
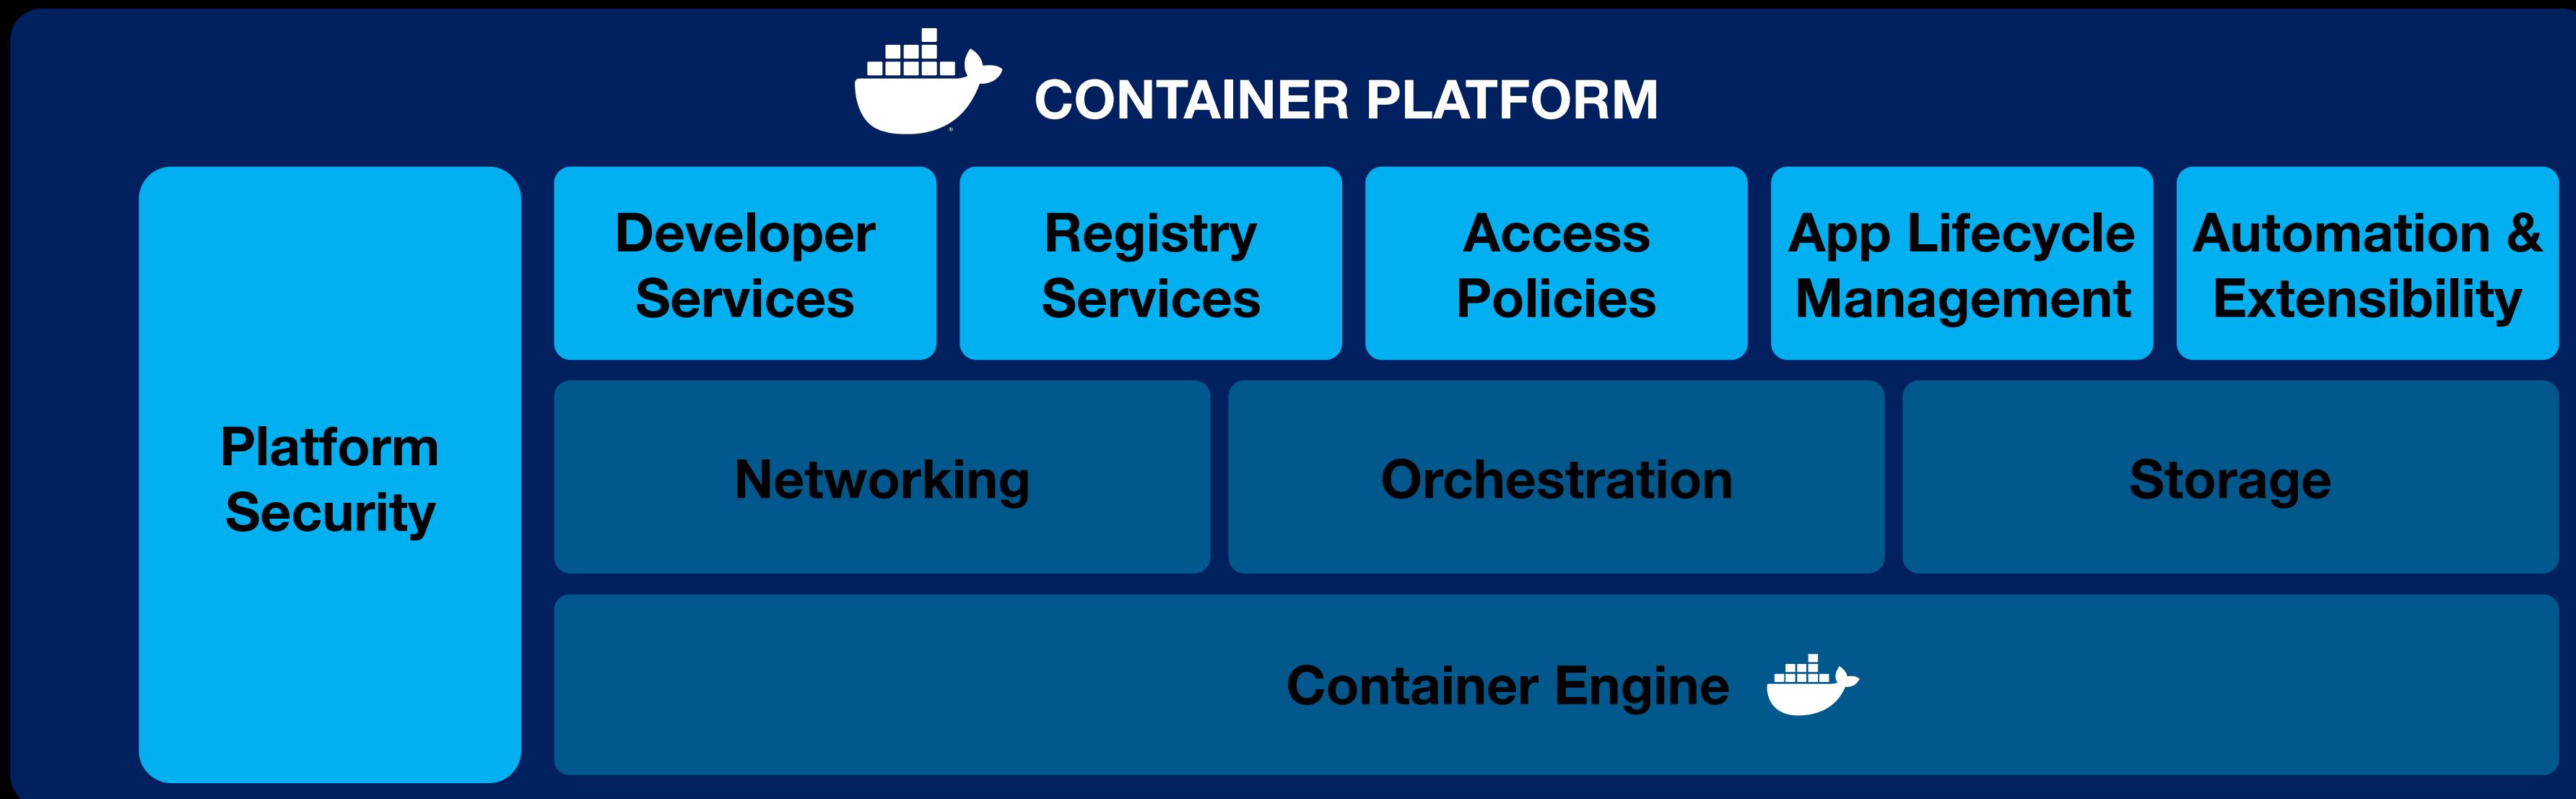
CPU

Memory

Liveness

File Descriptors

Storage Capacity



## 1. Development

Version control

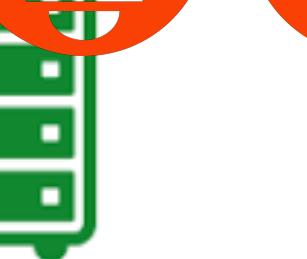
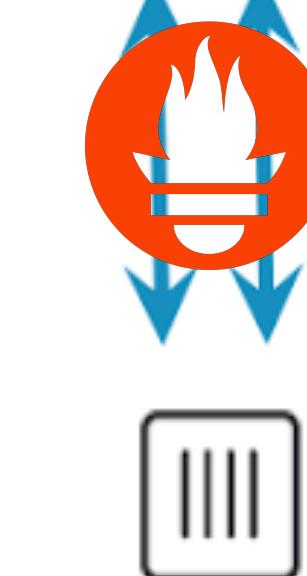
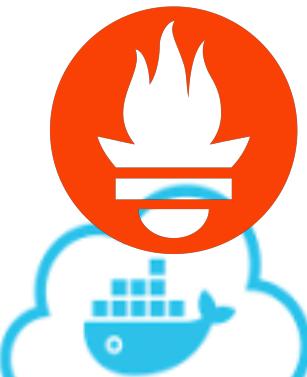


**Developer**



## 2. Test

Auto-builds



## 3. Stage / Production

Auto redeploy



**Sysadmin**



**QA / QE**

# Website Down?

- Total Downtime: Just under 4 minutes
- 502 error messages total: 12 000
- People affected by the 502 error who did not get their bargain: 400



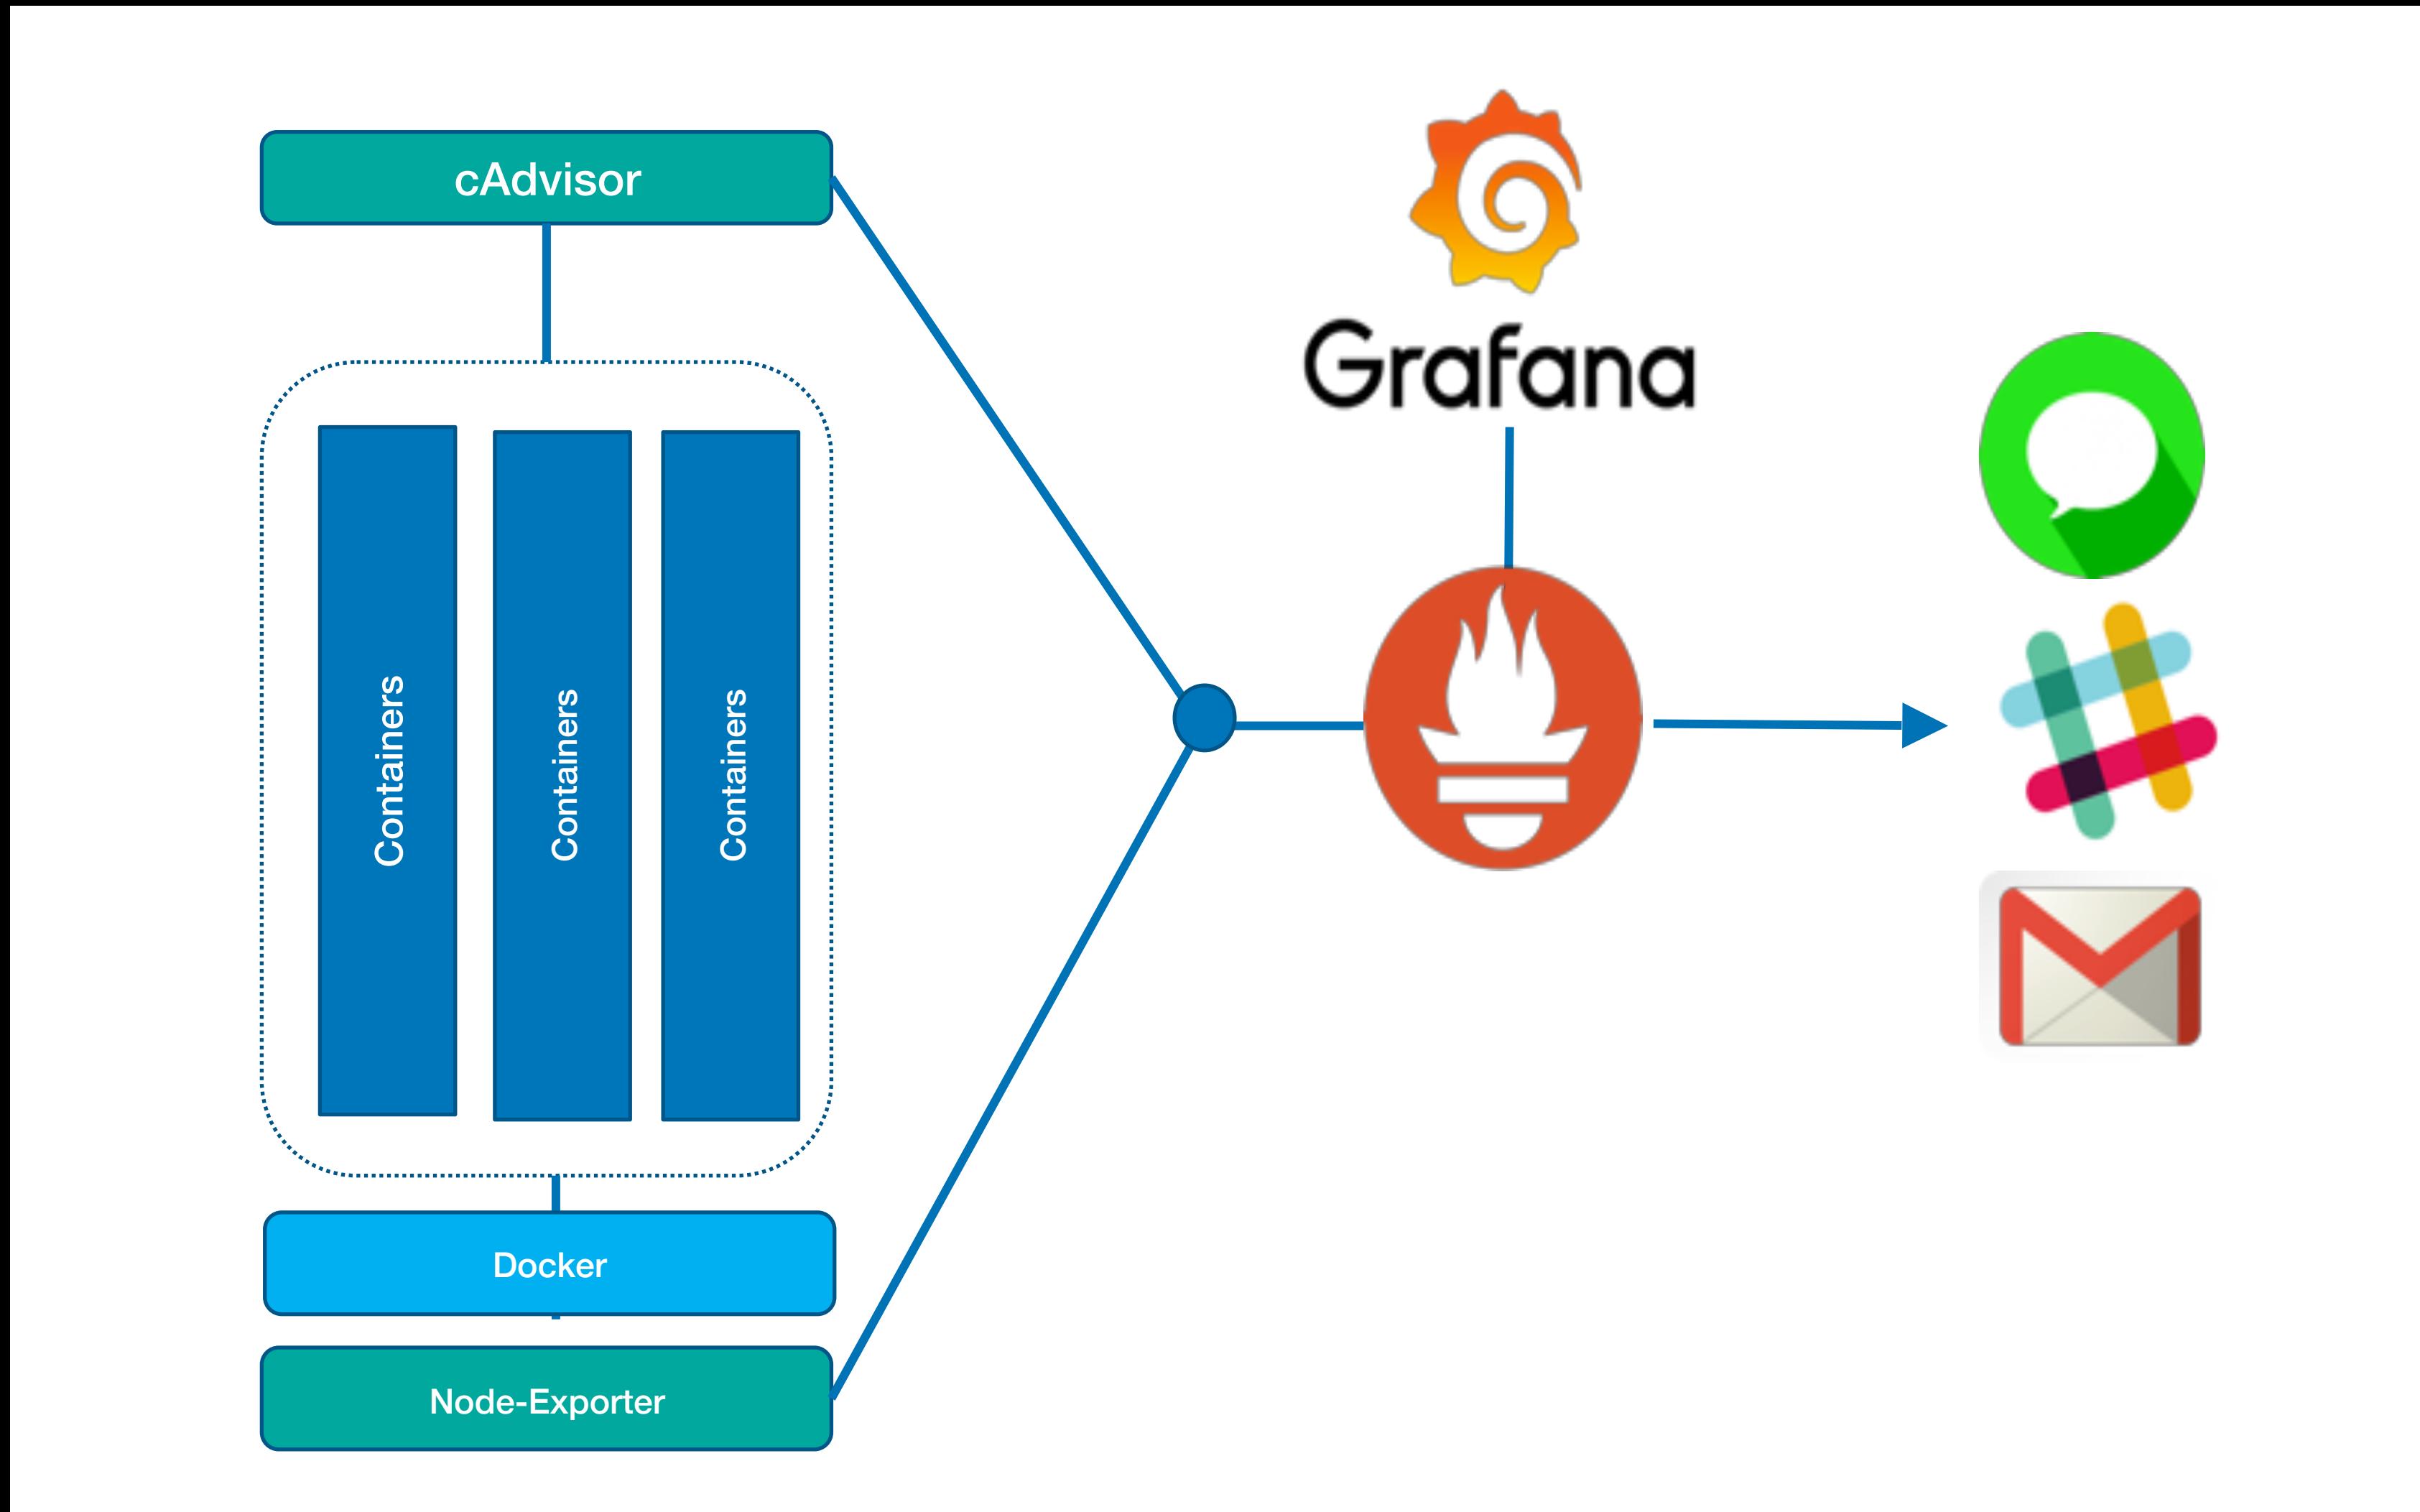


# 3 Easy Steps - Code

<https://about.gitlab.com>

- [ X ] Introduction
- [ X ] Operations Overview
- [ X ] What to Monitor
- [ ] Demo
- [ ] Best Practices & Recap

# DEMO



- [ X ] Introduction
- [ X ] Operations Overview
- [ X ] What to Monitor
- [ X ] Demo
- [ ] Best Practices & Recap

# Improve Incrementally



# Best Practices

- Start small & increment
- Don't Overlert yourself
- Set Resource Limits
- Aim for actionable Information
- Run separate from Workload
- Test for Failures
- Know your Failure Models





Kelsey Hightower

@kelseyhightower

Following

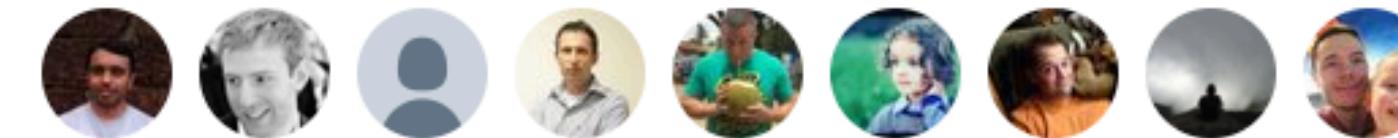


The thing you're trying to build should be more exciting than the tools used to build it.

4:56 PM - 11 Mar 2019

---

499 Retweets 1,841 Likes



---

36

499

1.8K



# Resources

- 56K.Cloud - <https://56K.Cloud>
- Prometheus - <https://github.com/vegasbrianc/prometheus>
- Monitoring Labs - [github.com/56kcloud/Training/](https://github.com/56kcloud/Training/)
- Docker Resource Link - <https://awesome-docker.netlify.com>
- GitLab Dashboards - <https://monitor.gitlab.net>

# Thank You

Brian Christner

[brian@56K.cloud](mailto:brian@56K.cloud)

@idomyowntricks